

---

# MACHINE LEARNING

## CHAPTER 0: INTRODUCTION

---

# Contact Information

---

**Instructor:** Qi Hao

**E-mail:** hao.q@sustc.edu.cn

**Office:** Nanshan iPark A7 Room 906

**Office Hours:** M 2:00-4:00pm

Available other times by appointment or the open door policy

**Office Phone:** (0755) 8801-8537

**QQ:** 1132148103 机器学习2020

**Web:** <http://hqlab.sustc.science/teaching/CS405>

**BB:** CS405

---

# Class Schedule

---

■ **Lectures:** W 08:00 am – 09:50 am Lychee Hills 1 Room 203

■ **Labs:** W 10:20 am – 12:10 pm Lychee Hills 6 Room 408  
W 16:20 am – 18:10 pm Lychee Hills 6 Room 406

■ **Grading policy:**

**Final Exam (in-class): 20%      Midterm Exam (take-home): 10%**

**Assignments (8~12 times): 20%      Quizzes (<=10 times): 10%**

**Lab Projects : 20%      Final Projects (4 per group): 20%**

**Bonus Credits: <5%**

90~93: A-

94~97: A

98~100: A+

80~82: B-

83~86: B

87~89: B+

70~72: C-

73~76: C

77~79: C+

60~62: D-

63~66: D

67~69: D+

# Textbook and Lecture Notes

---

## Textbooks:

- [1] Pattern Recognition and Machine Learning, by Christopher M. Bishop, 2006 Springer
- [2] Machine Learning in Action, by Peter Harrington, 2012, Manning

## Other books:

- [1] 机器学习, 周志华
- [2] Dive in Deep Learning, by Aston Zhang, Zachary C. Lipton, Mu Li, and Alexander J. Smola
- [3] Reinforcement Learning: An Introduction, by Richard S. Sutton
- [4] The Elements of Statistical Learning, by Trevor Hstie, Rober Tibshirani, Jerome Friedman

## Paper reading:

- [1] Ghahramani Z. Probabilistic machine learning and artificial intelligence, Nature, 2015
- [2] Lecun Y, Bengio Y, Hinton G. Deep learning, Nature, 2015
- [3] Littman M L. Reinforcement learning improves behavior from evaluative feedback, Nature, 2015

## Lecture notes:

<http://hqlab.sustc.science/teaching/CS405>

---

# Other Resources

---

**Assignment platform:** [bb.sustech.edu.cn](http://bb.sustech.edu.cn)

**Textbook resource:** <https://www.microsoft.com/en-us/research/people/cmbishop/#prml-book>

**Textbook Matlab codes :** <http://prml.github.io/>

**Matlab toolboxes :** Machine Learning, Neural Networks

---

# Teaching Objectives

---

- Fundamental knowledge about machine learning and pattern recognition, from Bayesian approaches to deep learning frameworks through lectures, quizzes and exercises
- Machine learning system development methods in Python based platforms (numpy, sciki-learn, tensorflow) through labs and projects
- Model-based and data-driven machine learning system design and integration skills through the final project, literature surveys and reports

# Lecture Schedule

---

Section 0	Course Introduction	
Section 1	Preliminary	(HW1)
Section 2	Probability Distributions	(HW2)
Section 3	Linear Regression and Classification	(HW3)
Section 4	Neural Networks	(HW4)
Section 5	Sparse Kernel Machine	(HW5)
Section 6	Clustering and EM learning	(HW6)
<i>-- Midterm Exam --</i>		
Section 7	Dimension Reduction and Feature Selection	(HW7)
Section 8	Ensemble Learning	(HW8)
Section 9	Bayesian Networks	(HW9)
Section 10	Hidden Markov Model	(HW10)
Section 11	Markov Decision Process	(HW11)
Section 12	Reinforcement Learning	(HW12)
<i>-- Final Exam --</i>		

---

# Lab Schedule

---

Section 0    Lab Introduction

Section 1    Preliminary

Section 2    Bayes

Section 3    Regression and Classification

--*Final Project Proposal*--

Section 4    Decision Tree

Section 5    Random Forest (Ensemble Learning)

Section 6    KNN and Support Vector Machine

Section 7    K-Mean and EM Clustering

Section 8    Neural Network (I)

Section 9    Neural Network (II)

Section 10    Neural Network (III)

Section 11    Reinforcement Learning

--*Final Project Report*--

---

# Final Projects

---

- [1] Reinforcement learning based planning using a self-driving car simulator
  - [2] Generation of annotated self-driving datasets with the CARLA simulator
  - [3] Segmentation of 2D/3D measurements for self-driving applications
  - [4] Detection and recognition of traffic signs for self-driving applications
  - [5] Detection and tracking of 2D/3D objects for self-driving applications
  - [6] Federated learning for model fusion of networked vehicle applications
-

# Bonus Credits

---

- AI companies
- Survey Papers
- Attendance
- Bonus Credits



# Plagiarism

---

**From Spring 2018**, the plagiarism policy applied by the Computer Science and Engineering department is the following:

- \* If an assignment is found to be plagiarized, the first time the score of the assignment will be 0.
- The second time the score of the course will be 0.

As it may be difficult when two assignments are identical or nearly identical who actually wrote it, the policy will apply to BOTH students, unless one confesses having copied without the knowledge of the other.

---

# What is OK, and what isn't OK

---

## It's OK

- to work on an assignment with a friend, and think together about the program structure, share ideas and even the global logic. At the time of actually writing the code, you should write it alone.
- to use in an assignment a piece of code found on the web, as long as you indicate in a comment where it was found and don't claim it as your own work.
- to help friends debug their programs (you'll probably learn a lot yourself by doing so).
- to show your code to friends to explain the logic, as long as the friends write their code on their own later.

## It's NOT OK

- **to take the code of a friend, make a few cosmetic changes (comments, some variable names) and pass it as your own work.**
-

# **Make a Promise to Keep**

---

**Sign**

the “Student Commitment for Assignments”

**Keep**

the promise during the whole semester!

---

# Outlines

---

- Framework
  - Problem Statement
  - Related Areas
  - History
  - Datasets and Learning Models
  - Optimization Methods
  - Algorithms
  - Examples
-

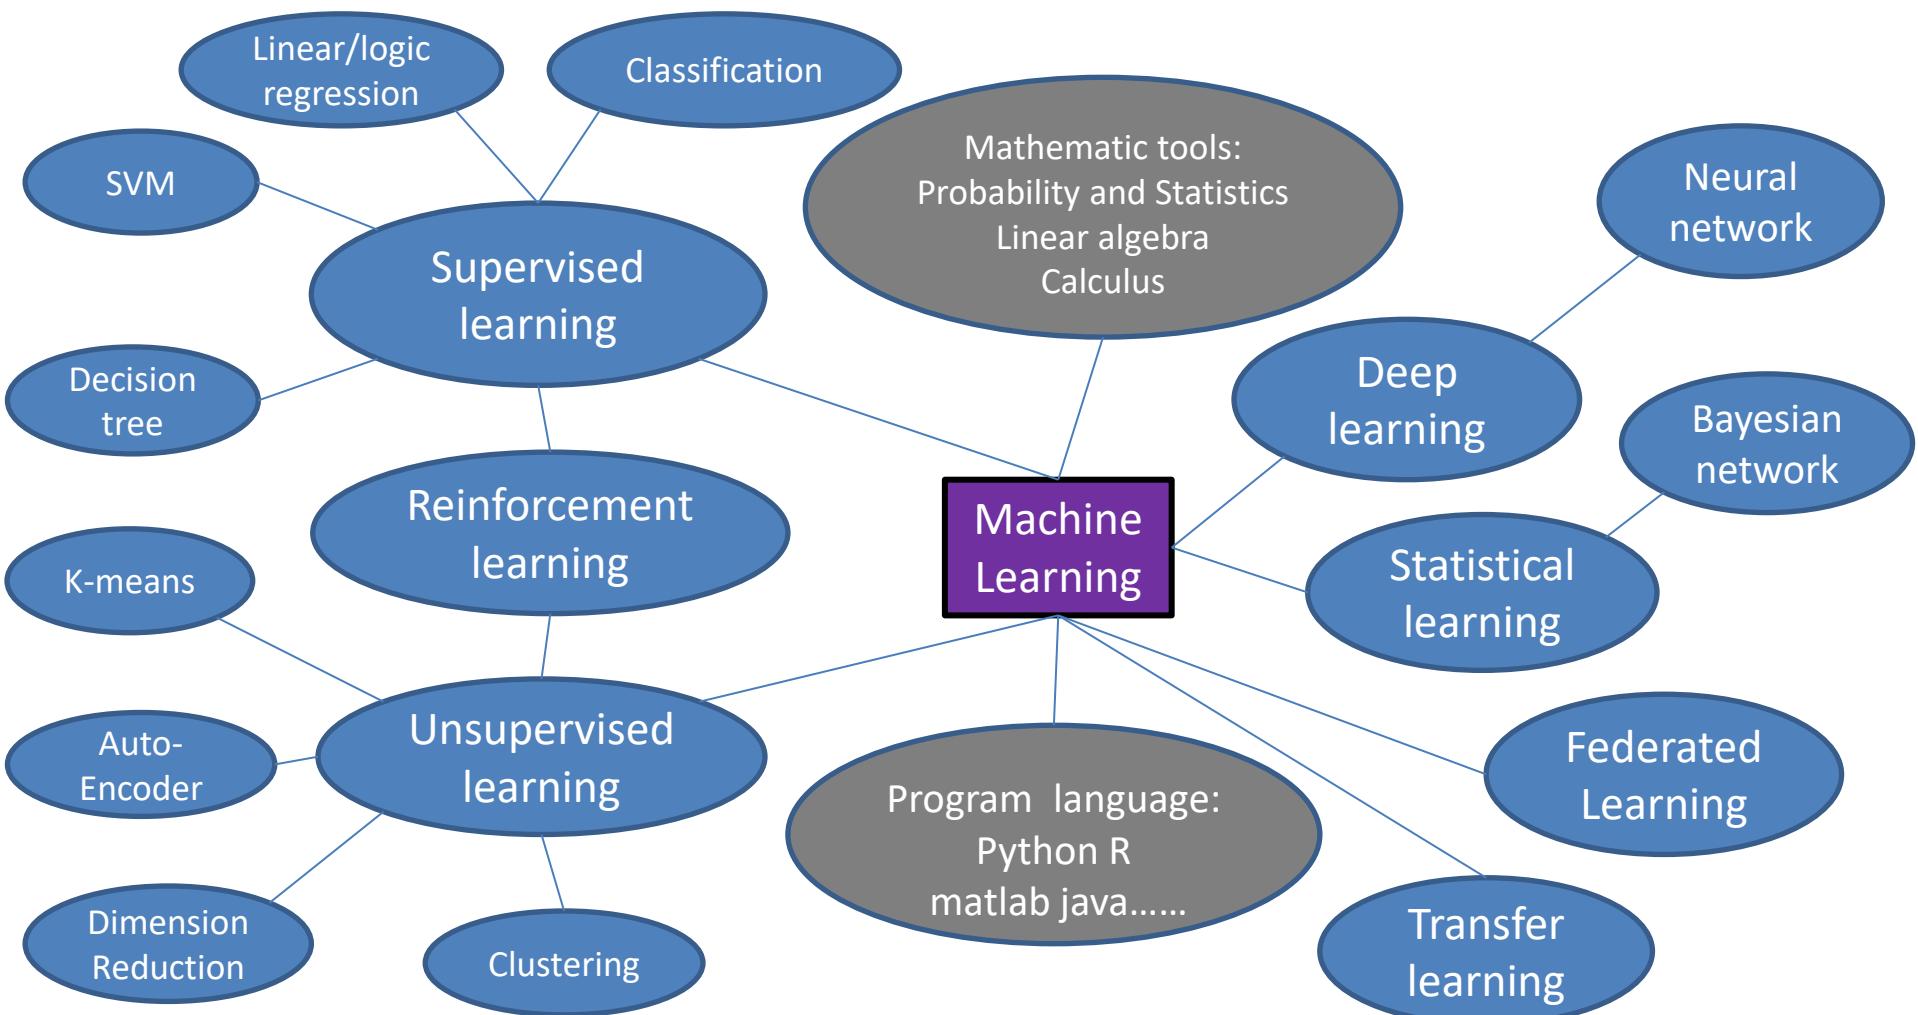
# Outlines

---

- Framework
  - Problem Statement
  - Related Areas
  - History
  - Datasets and Learning Models
  - Optimization Methods
  - Algorithms
  - Examples
-

# Framework

---



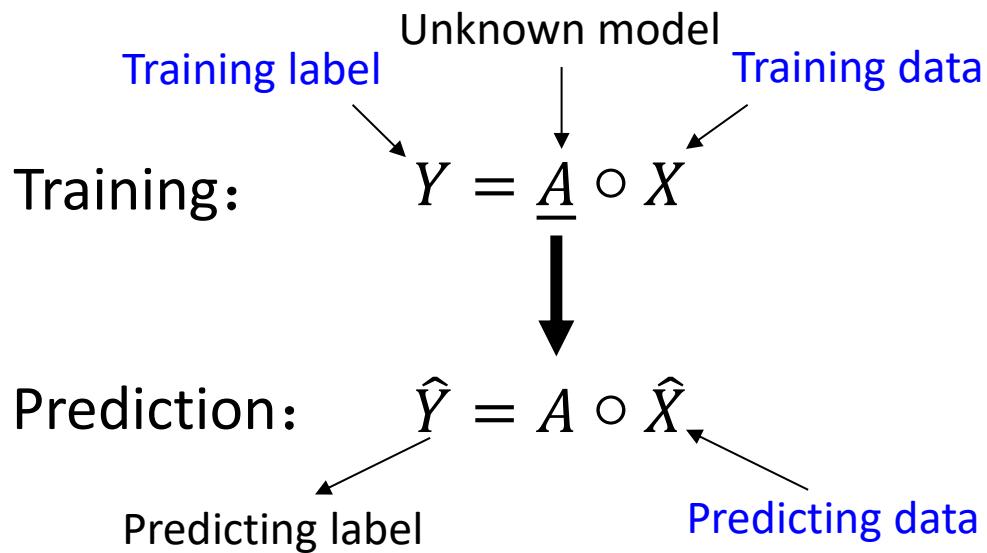
# Outlines

---

- Framework
  - Problem Statement
  - Related Areas
  - History
  - Datasets and Learning Models
  - Optimization Methods
  - Algorithms
  - Examples
-

# Problem Statement

- Problem: Predict the label  $\hat{Y}$  and data  $\hat{X}$  with training set  $(X, Y)$  ?

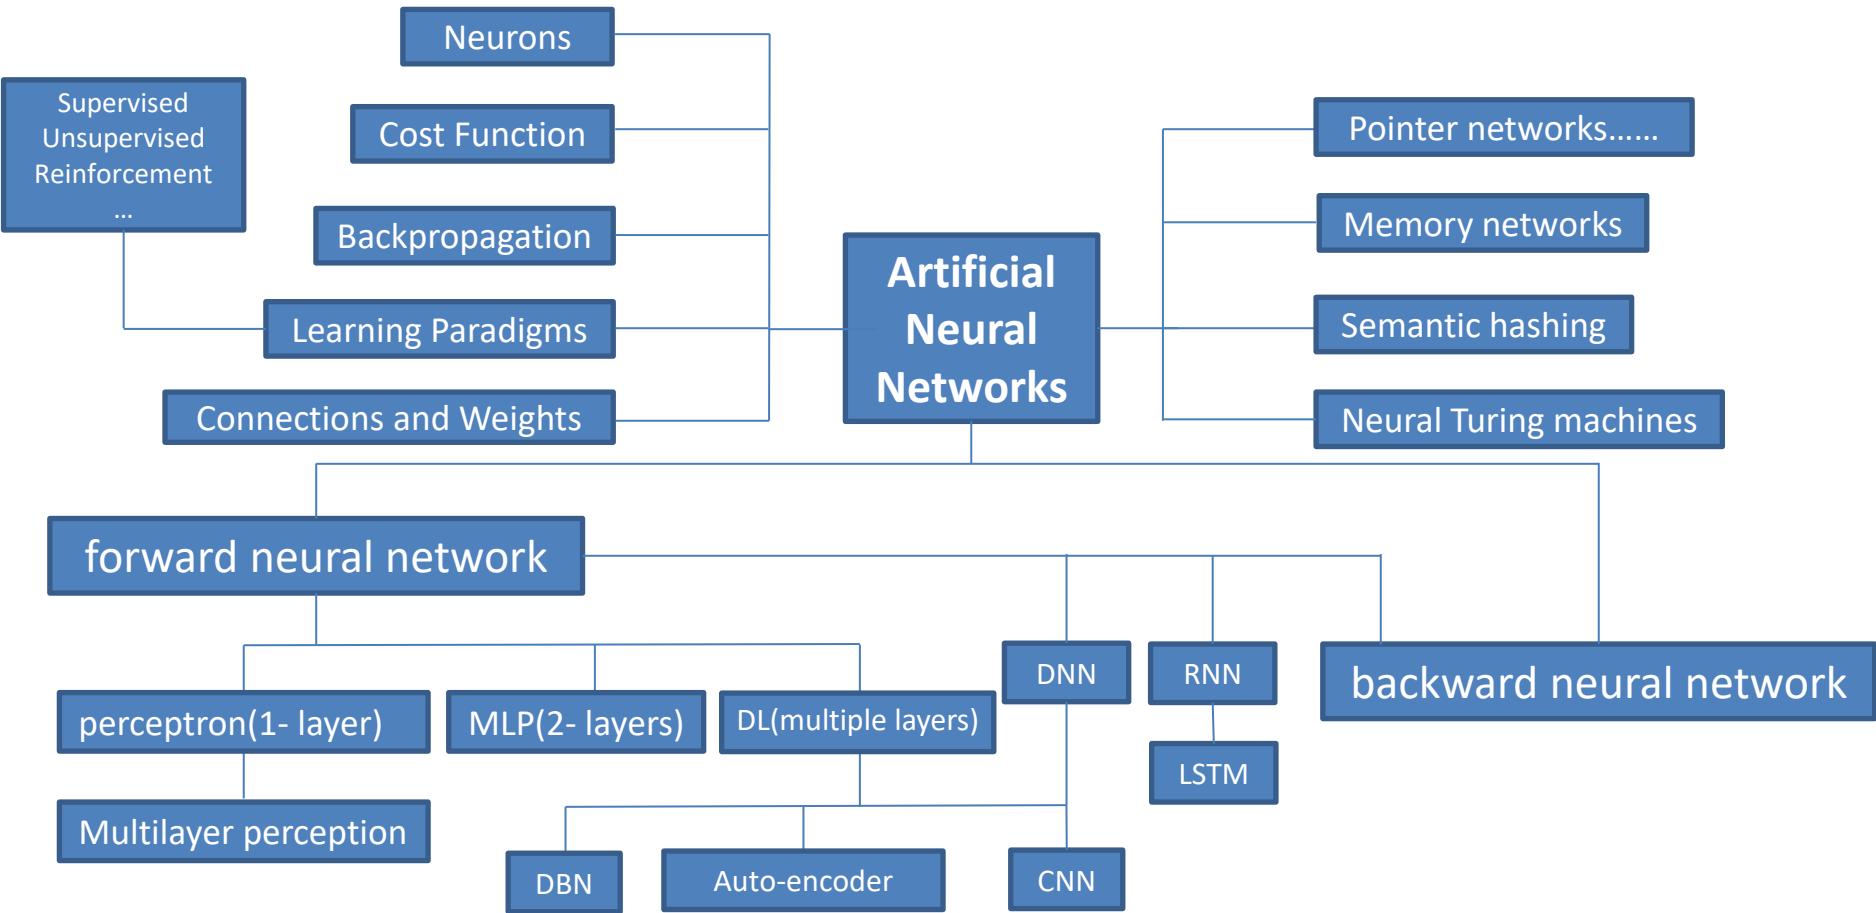


$\begin{cases} Y \text{ and } X \text{ is known: supervised learning} \\ Y \text{ or } X \text{ is unknown: unsupervised learning} \end{cases}$      $\begin{cases} Y, \hat{Y} \text{ are continuous: Regression} \\ Y, \hat{Y} \text{ are discrete: classification} \end{cases}$

$Y$  is known and  $\text{Dim}(Y) > \text{Dim}(X)$  : dimensionality reduction

# Neural Network Models

---



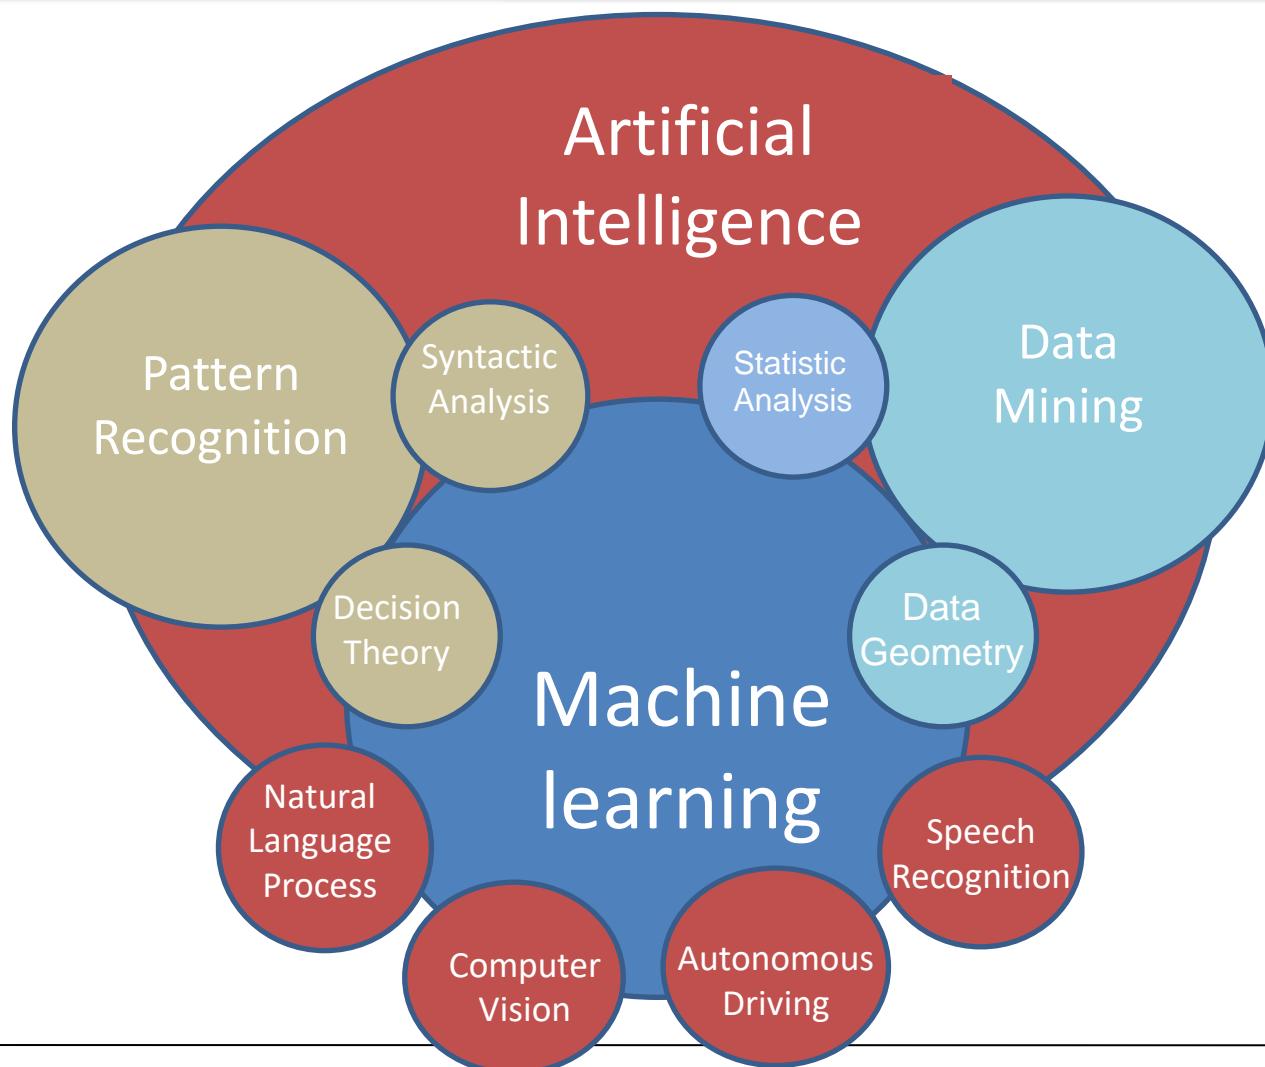
# Outlines

---

- Framework
  - Problem Statement
  - Related Areas
  - History
  - Datasets and Learning Models
  - Optimization Methods
  - Algorithms
  - Examples
-

# The Whole Picture

---



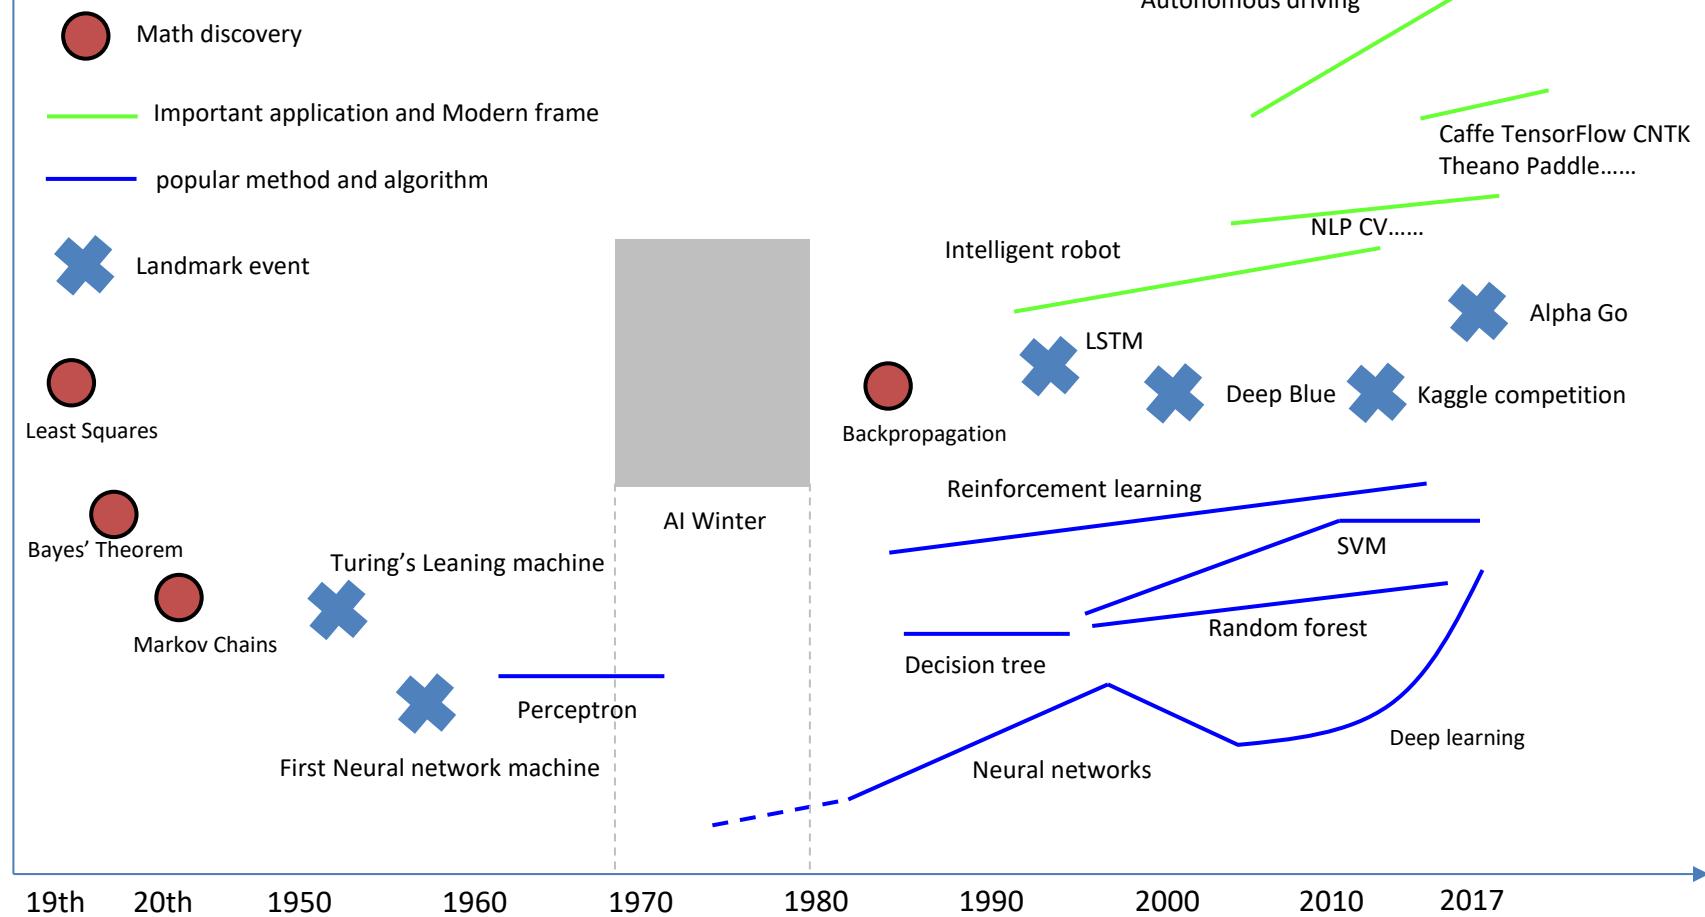
# Outlines

---

- Framework
  - Problem Statement
  - Related Areas
  - History
  - Datasets and Learning Models
  - Optimization Methods
  - Algorithms
  - Examples
-

# History

## Popularity and degree of application



# Outlines

---

- Framework
  - Problem Statement
  - Related Areas
  - History
  - Datasets and Learning Models
  - Optimization Methods
  - Algorithms
  - Examples
-

# Machine learning

---

- **Machine Learning**—minimization of some loss function for generalizing data sets with models.
  
- **Datasets** —annotated, indexed, organized
  
- **Models** —tree, distance, probabilistic, graph, bio-inspired
  
- **Optimization** —algorithms can minimize the loss.

# Datasets

---

- Collection
  - Storage
  - Annotation
  - Indexing
  - Organization
  - Access
-

# Simulators

---

- Data visualization
- Generate training data
- Algorithm evaluation

# Benchmark Metrics

---

- System functionalities
- System scalability
- System robustness
- System efficiency

# Models

---

- Tree Models
- Distance-based Models
- Probabilistic Models
- Neural Network Models
- Graph-based Models

# Models

---

- Boosting
  - Mixed Models
  - Ensemble Learning
-

# Outlines

---

- Framework
  - Problem Statement
  - Related Areas
  - History
  - Datasets and Learning Models
  - Optimization Methods
  - Algorithms
  - Examples
-

# Machine learning and Optimization

---

- **Machine Learning**—minimization of some loss function for generalizing data sets with models.
  
- **Datasets** —annotated, indexed, organized
  
- **Models** —tree, distance, probabilistic, graph, bio-inspired
  
- **Optimization** —algorithms can minimize the loss.

# What is optimization?

---

- Finding (one or more) minimizer of a function subject to constraints

$$\arg \min_x f_0(x)$$

$$s.t. f_i(x) \leq 0, i = \{1, \dots, k\}$$

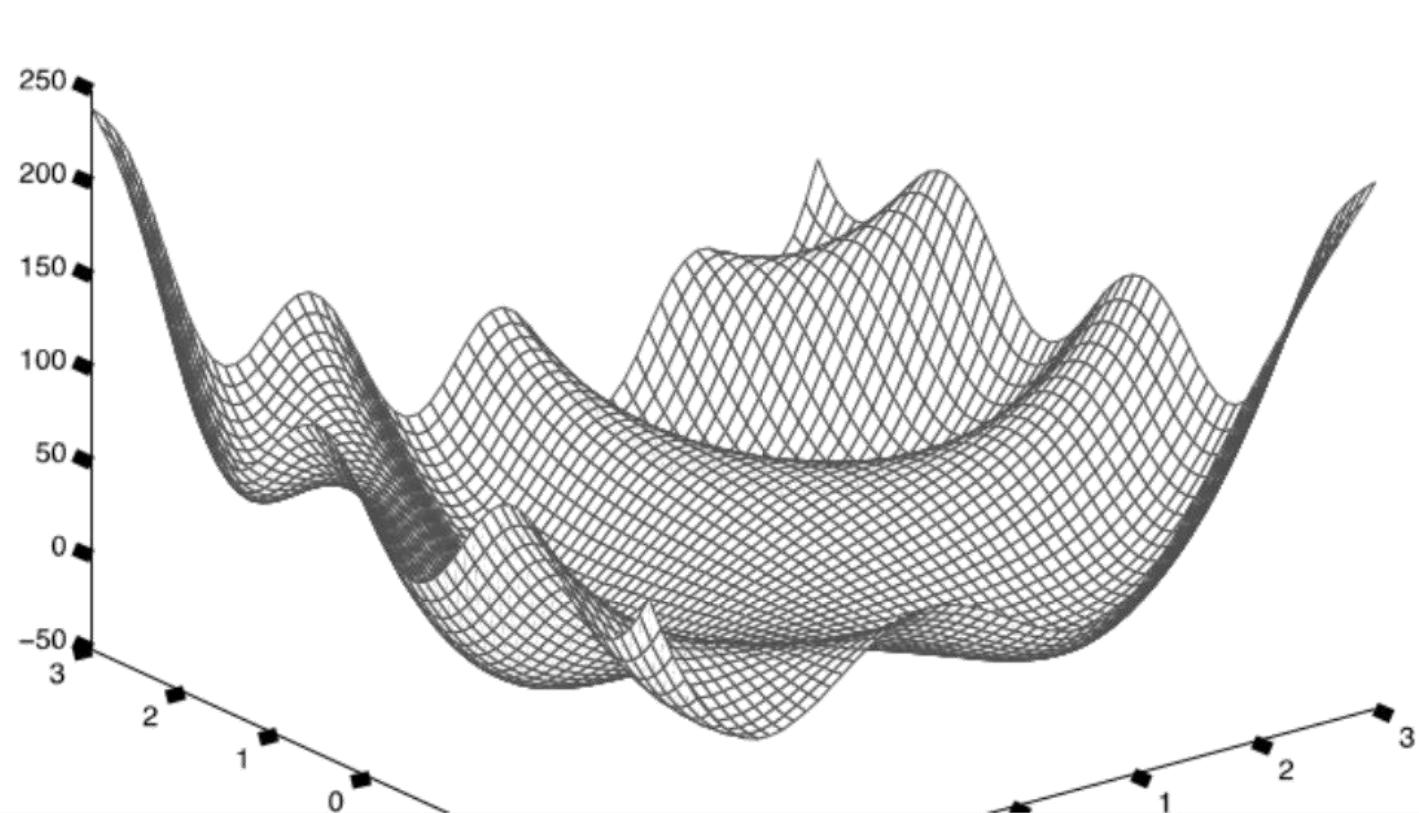
$$s.t. h_j(x) = 0, j = \{1, \dots, l\}$$

- Most of the machine learning problems are, in the end, optimization problems
-

# General Problem

---

■ Minimize  $f(x)$



# Linear Optimization

---

$$Y = AX + w$$

$$w \sim \mathcal{N}(0, R)$$

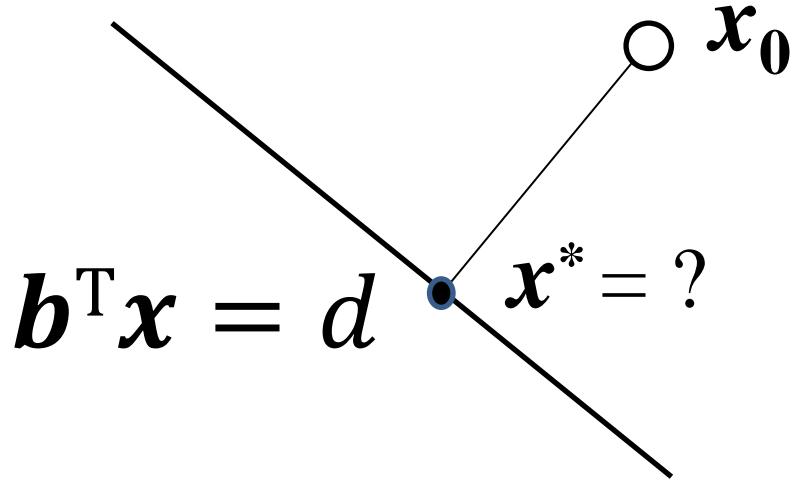
$$X^* = \min_X (Y - AX)^T R^{-1} (Y - AX)$$

$$\frac{\partial}{\partial X^T} (Y - AX)^T R^{-1} (Y - AX) = 0$$

$$\Rightarrow X^* = (A^T R^{-1} A)^{-1} A^T R^{-1} Y$$

# Linear Optimization

---



$$\mathbf{x}^* = \mathbf{x}_0 - \frac{(\mathbf{b}^T \mathbf{x}_0 - d)\mathbf{b}}{\mathbf{b}^T \mathbf{b}}$$

$$\mathbf{x}^* = \min_{\mathbf{x}} (\mathbf{x} - \mathbf{x}_0)^T (\mathbf{x} - \mathbf{x}_0)$$

$$\text{s. t. } \mathbf{b}^T \mathbf{x} - d = 0$$

---

# Nonlinear Optimization

---

## ■ Convex Optimization

- Unconstrained optimization
- Constrained optimization
- SVMs and Bayesian models

## ■ Non-convex Optimization

- Heuristic algorithms
  - Random search
-

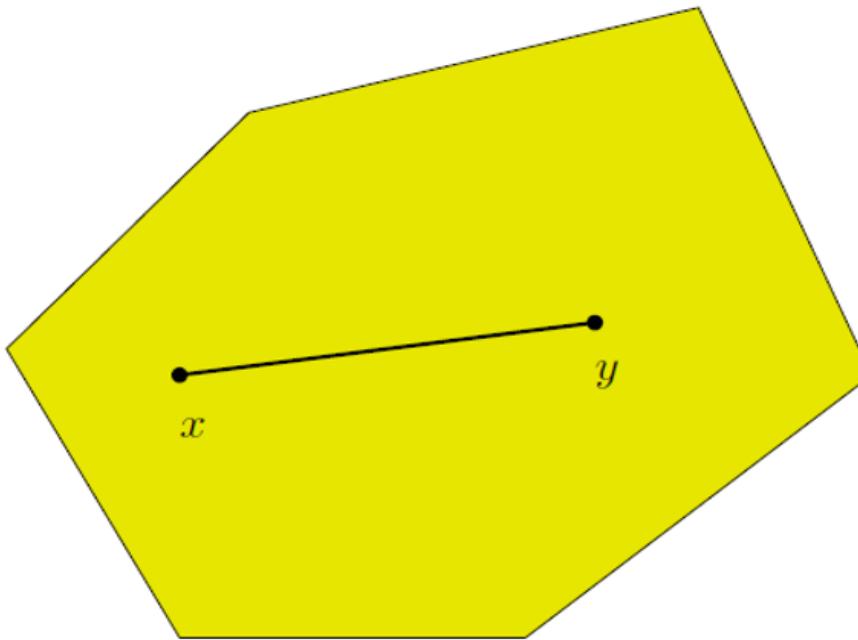
# What is Convex?

---

## ■ Convex sets

Def: A set  $C \subseteq \mathbb{R}$  is convex if for  $x, y \in C$ ;  $a \in [0, 1]$

$$ax + (1 - a)y \in C$$

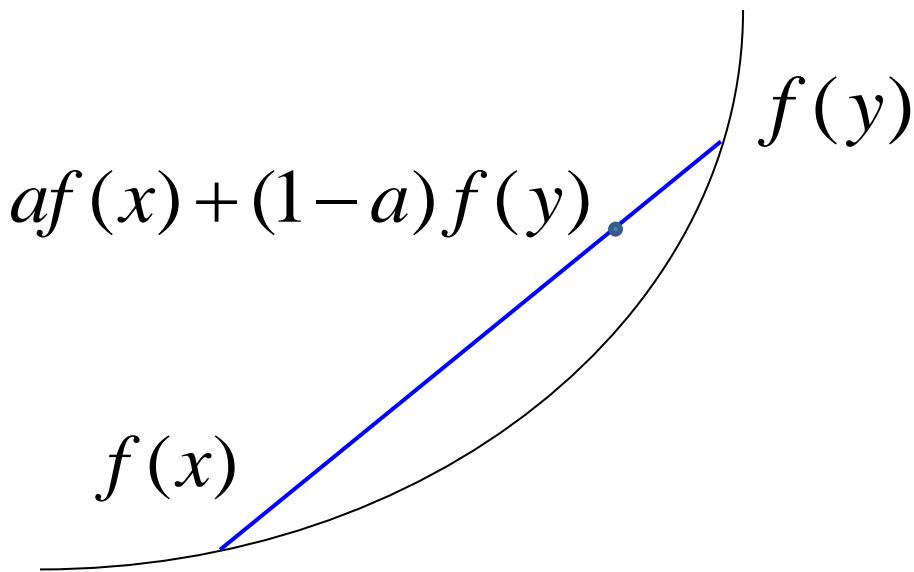


# What is Convex?

---

## ■ Convex functions

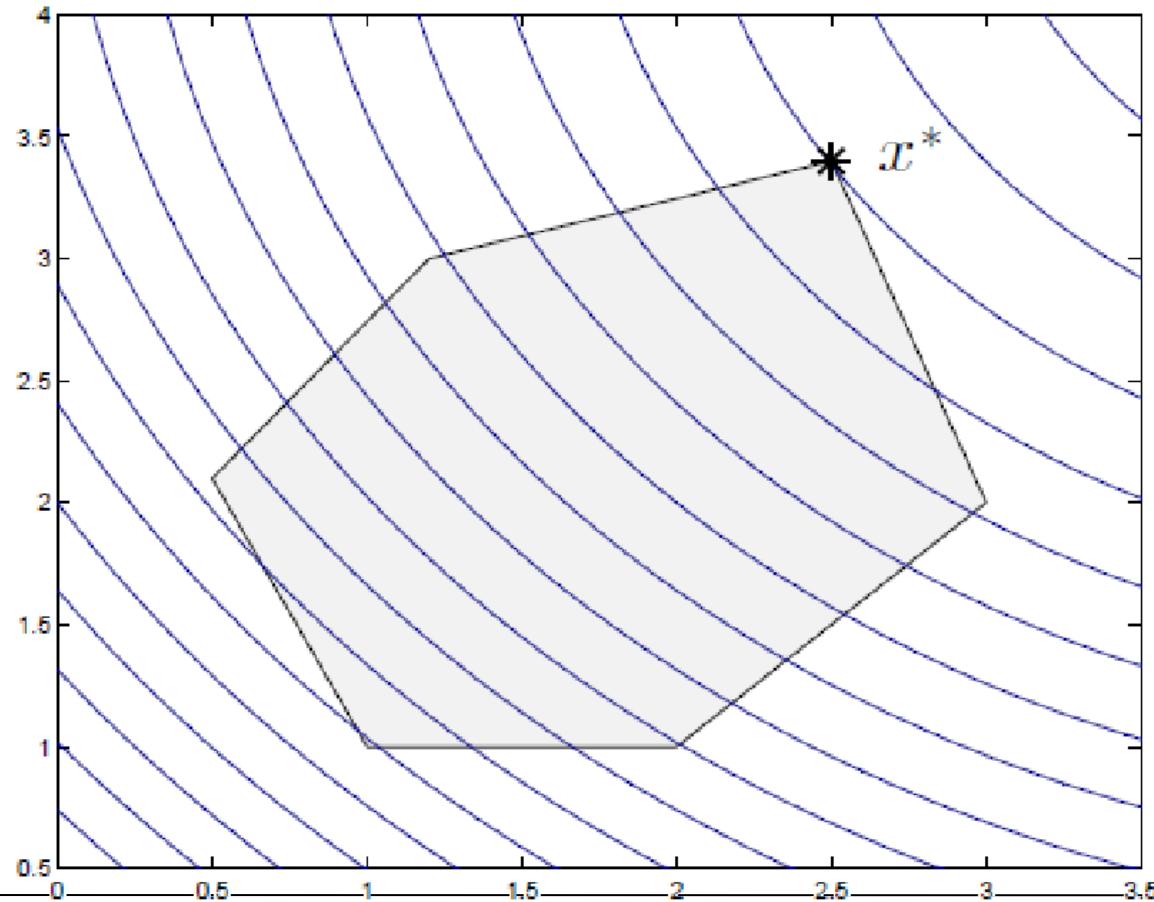
$$f(ax + (1-a)y) \leq af(x) + (1-a)f(y)$$



# Convex Optimization

---

- Local minimizer = Global minimizer



# Convex Optimization

---

- Unconstrained optimization
  - Gradient descent
  - Gauss-Newton's method
  - Batch learning
  - Stochastic Gradient Descent
  
- Constrained optimization
  - Lagrange methods
  - Bayesian methods

# Convex optimization

---

## ■ Unconstrained optimization

- Gradient descent
- Gauss-Newton's method
- Batch learning
- Stochastic Gradient Descent

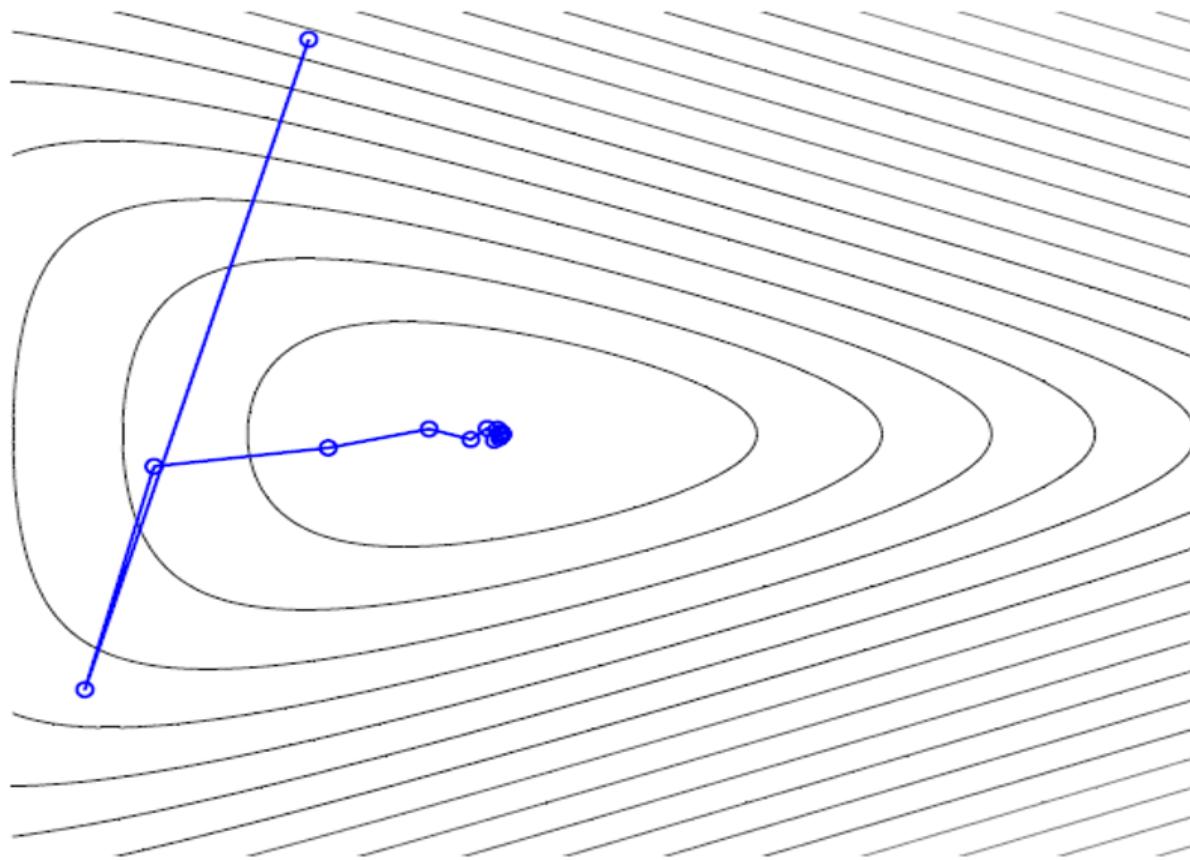
## ■ Constrained optimization

- Lagrange methods
  - Bayesian methods
-

# Gradient Descent

---

$$f(x_{t+1}) = f(x_t) - \eta \nabla f(x_t)^T (x - x_t)$$



# Gauss-Newton's Method

---

- Idea: use a second-order approximation to function

$$f(x + \Delta x) \approx f(x) + \nabla f(x)^T \Delta x + \frac{1}{2} \Delta x^T \nabla^2 f(x) \Delta x$$

- Choose  $\Delta x$  to minimize above:

$$\Delta x = -[\nabla^2 f(x)]^{-1} \nabla f(x)$$

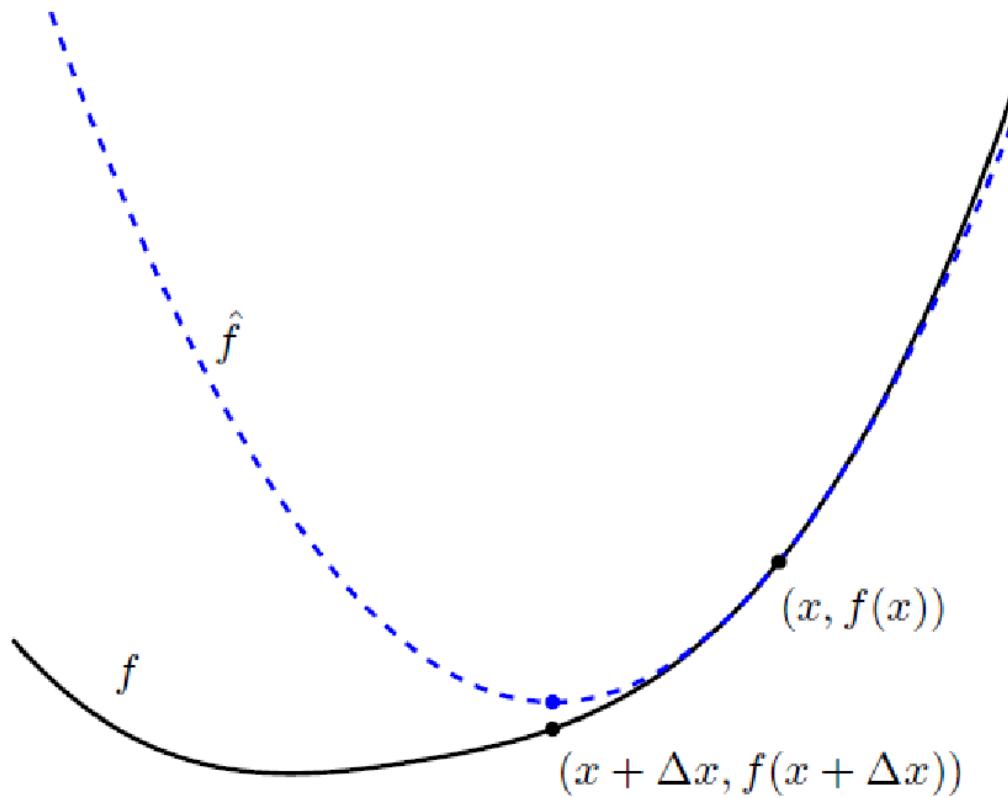
- This is descent direction:

$$\nabla f(x)^T \Delta x = -\nabla f(x)^T [\nabla^2 f(x)]^{-1} \nabla f(x) < 0$$

# Gauss-Newton's Method

---

$\hat{f}$  is 2-order approximation,  $f$  is true function.



# Batch Gradient Descent

---

- Minimize empirical loss, assuming it's convex and unconstrained

- Gradient descent on the empirical loss
  - At each step:

$$w^{(k+1)} \leftarrow w^{(k)} - \eta_t \left( \frac{1}{n} \sum_{i=1}^n \frac{\partial L(w, x_i, y_i)}{\partial w} \right)$$

- Note: at each step, gradient is the average of the gradient for all samples ( $i=1, \dots, n$ )
  - Very slow when  $n$  is very large

# Stochastic Gradient Descent

---

- Alternative: compute gradient from just one (or a few samples)
- Known as stochastic gradient descent:
  - At each step,

$$w^{(k+1)} \leftarrow w^{(k)} - \eta_t \frac{\partial L(w, x_i, y_i)}{\partial w}$$

(choose one sample  $i$  and compute gradient for that sample only)

---

# Convex Optimization

---

## ■ Unconstrained optimization

- Gradient descent
- Gauss-Newton's method
- Batch learning
- Stochastic Gradient Descent

## ■ Constrained optimization

- Lagrange methods
  - Bayesian methods
-

# Lagrange Methods

---

- Start with optimization problem:

$$\arg \min_x f_0(x)$$

$$s.t. f_i(x) \leq 0, i = \{1, \dots, k\}$$

$$s.t. h_j(x) = 0, j = \{1, \dots, l\}$$

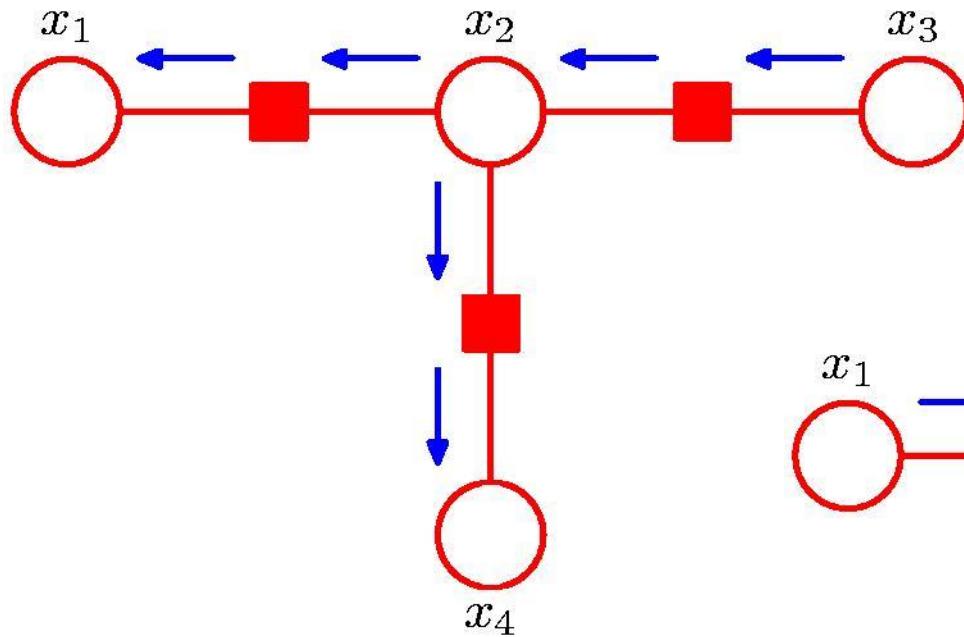
- Is equivalent to min-max optimization:

$$\arg \min_x \left[ \max_{\lambda \geq 0, \gamma > 0} \left( f_0(x) + \sum_{i=1}^k \lambda_i f_i(x) + \sum_{j=1}^l \gamma_j h_j(x) \right) \right]$$

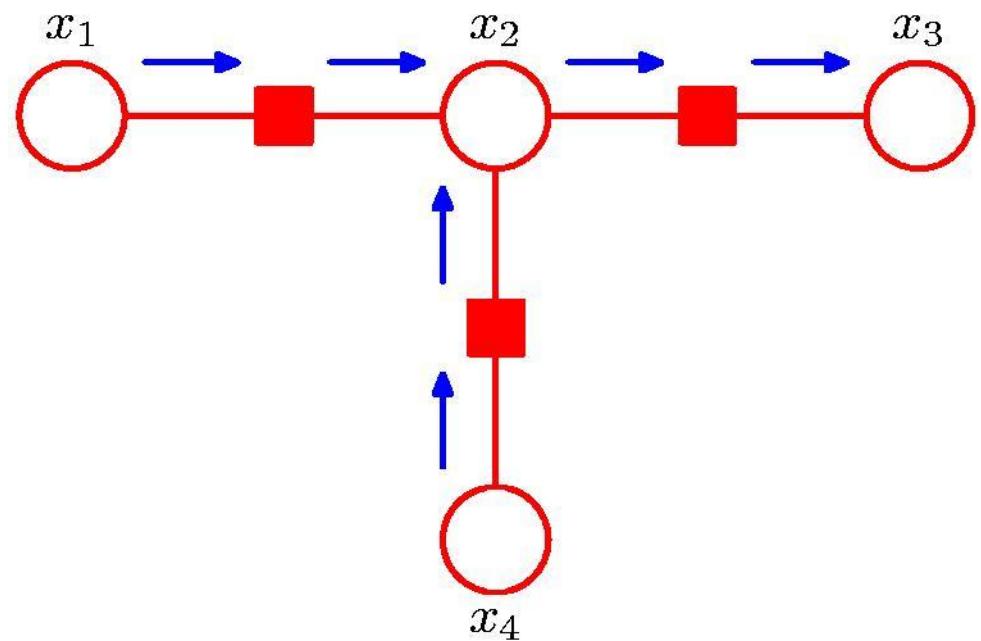
---

# Bayesian Methods

---



Random variable  
 factor



# Convex Optimization for Machine Learning

---

- Support Vector Machine

- Lagrange methods

- Bayesian Models:

- Expectation-Maximization methods
  - Variational methods
  - Graph optimization

# Non-convex Optimization

---

## ■ Convex Optimization

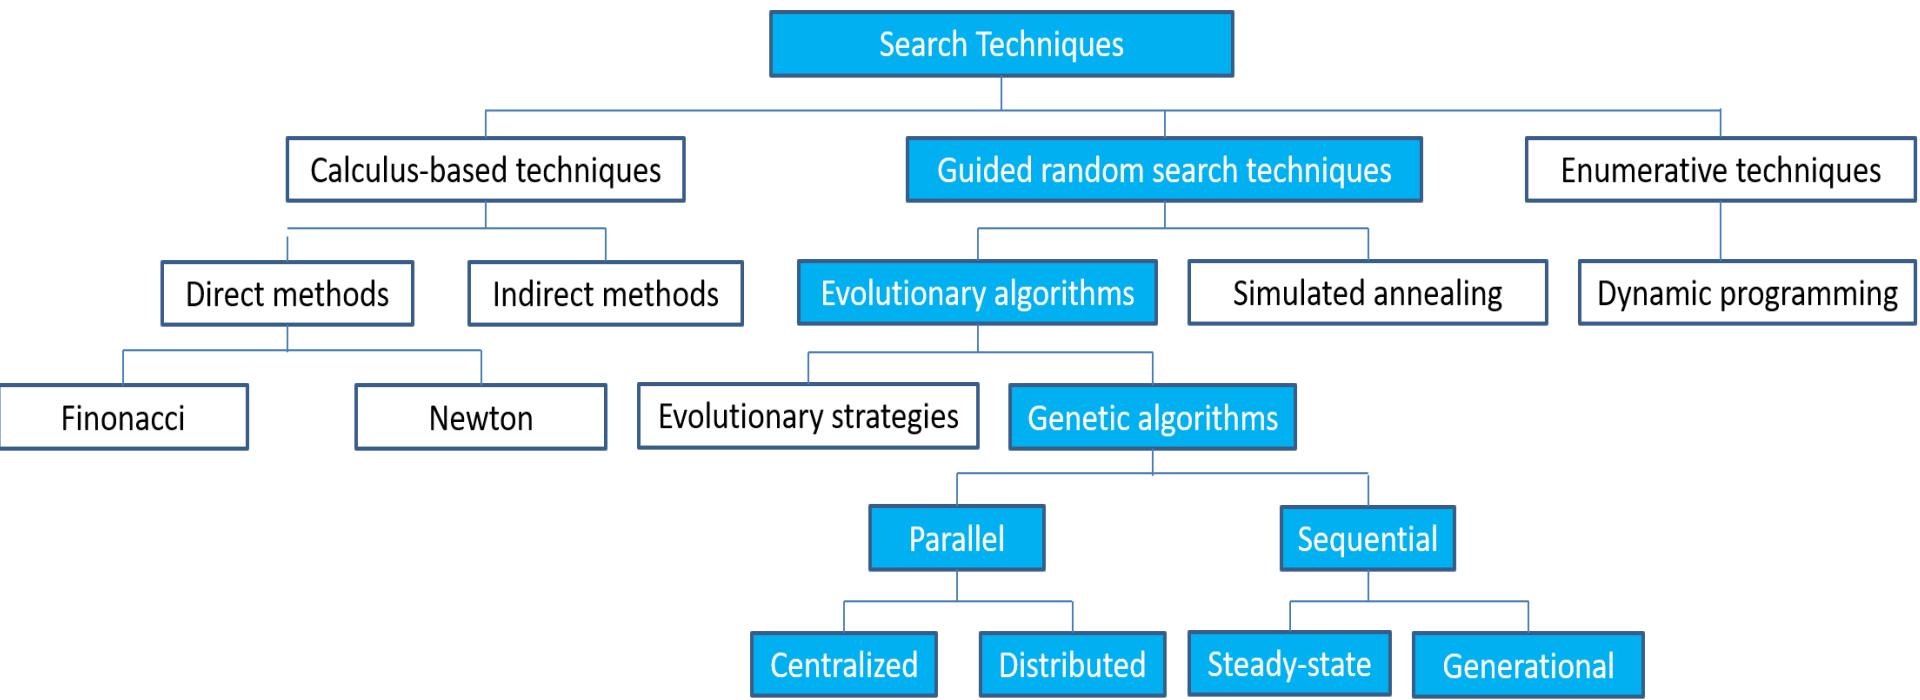
- Unconstrained optimization
- Constrained optimization

## ■ Non-convex Optimization

- Heuristic algorithms
  - Random search
-

# Heuristic and Random Search

---



# Outlines

---

- Framework
  - Problem Statement
  - Related Areas
  - History
  - Datasets and Learning Models
  - Optimization Methods
  - Algorithms
  - Examples
-

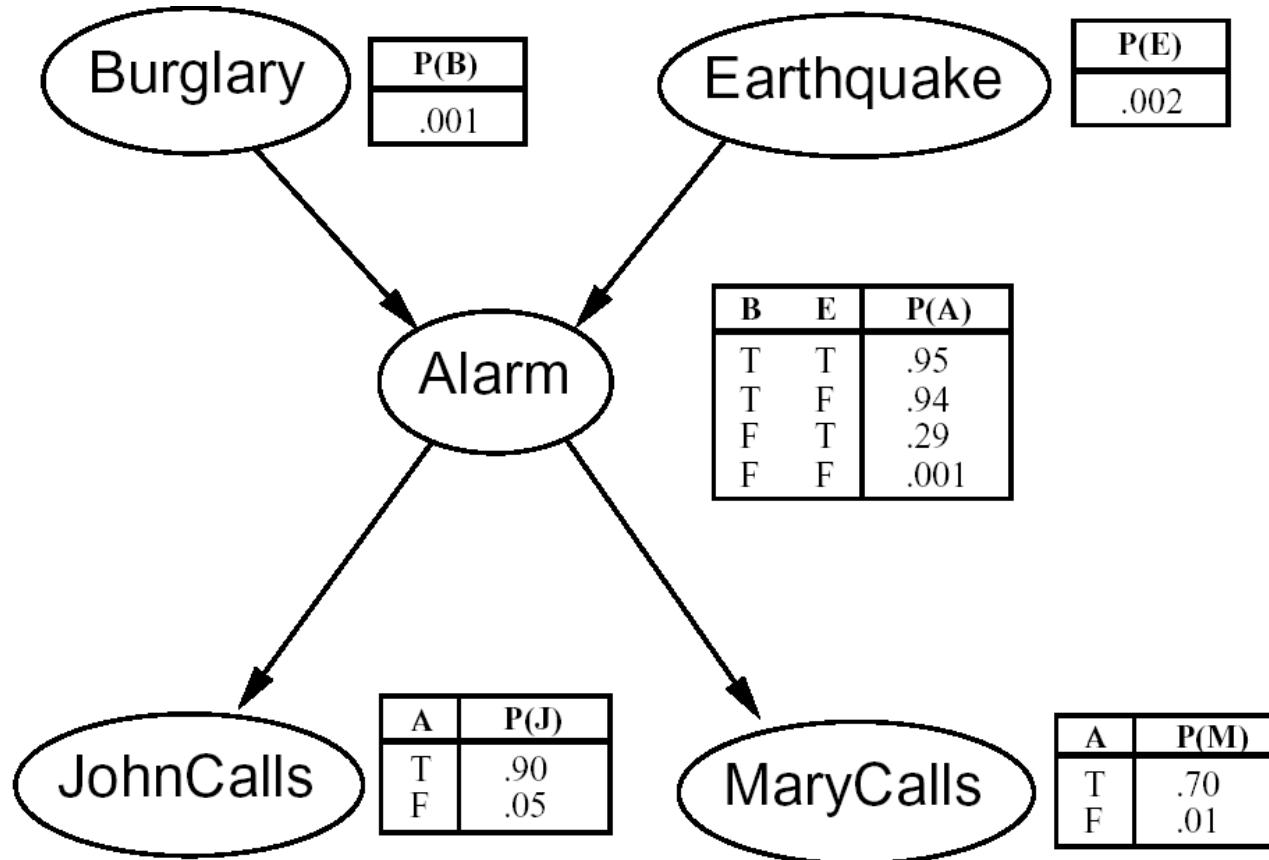
# Algorithms

---

- Bayes
  - KNN and K-means
  - Decision tree
  - Support Vector Machine
  - Boosting and Ensemble Learning
  - Linear Statistical Learning (PCA, ICA, NMF)
  - Nonlinear Statistical Learning (Manifold learning)
  - Deep Neural Networks
  - Generative Adversarial Networks
  - Bayesian Networks
  - Reinforcement Learning
  - Federated Learning
-

# Bayes

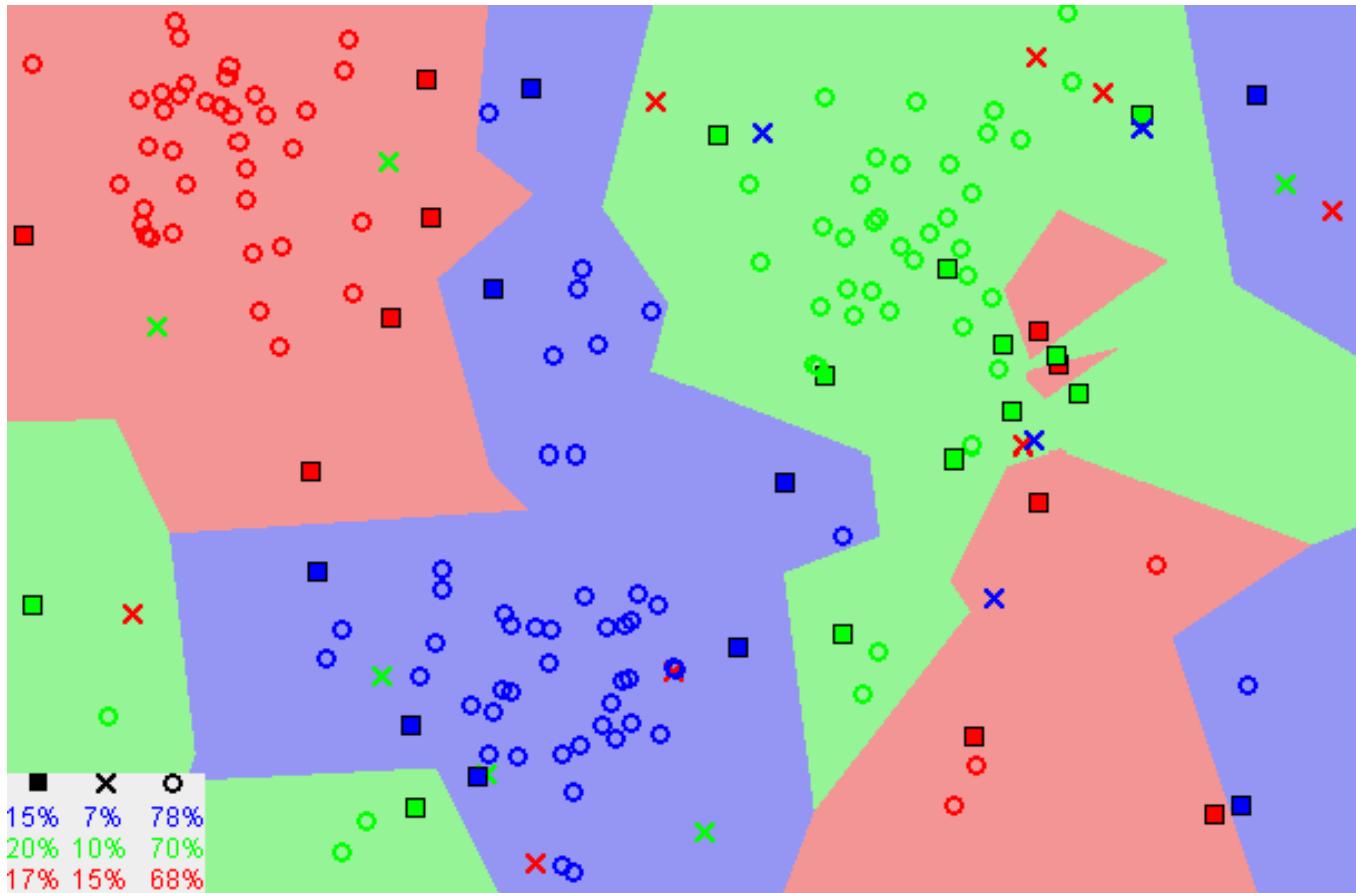
---



# K-Nearest Neighbors

---

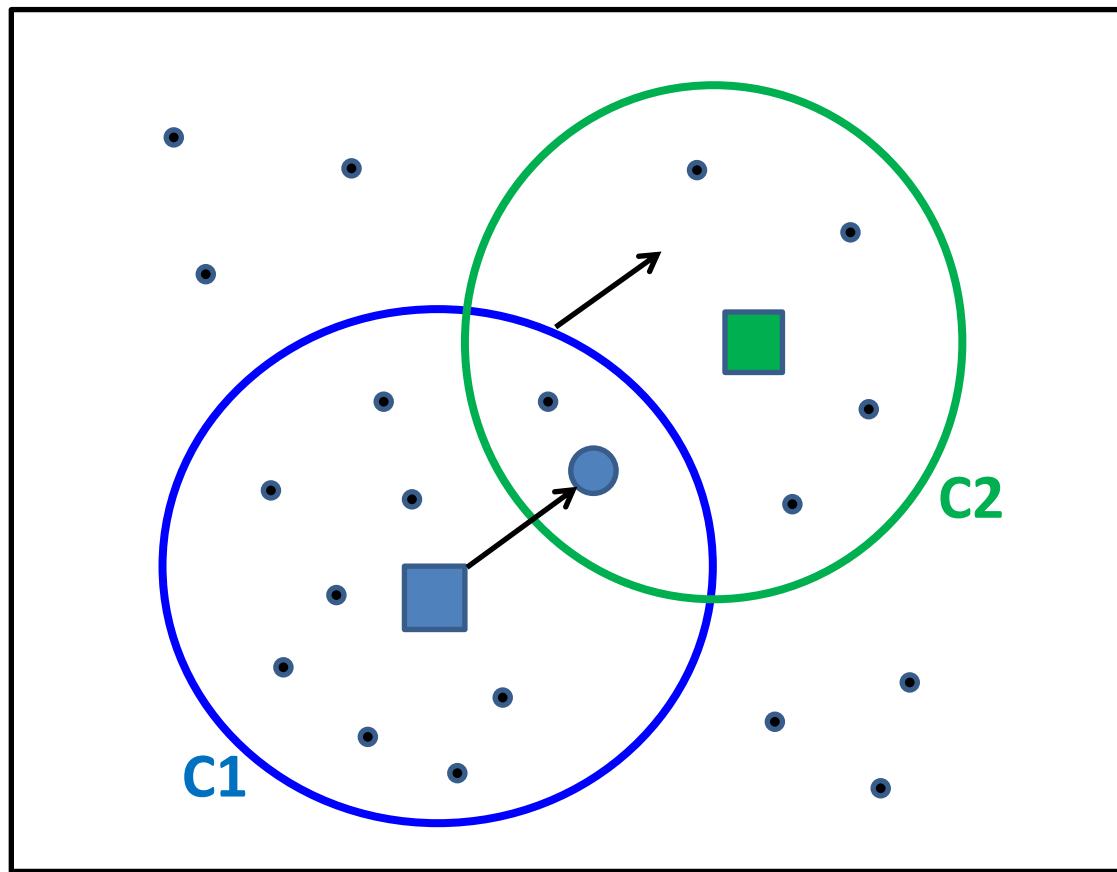
- Use training data for classification



# K-Means

---

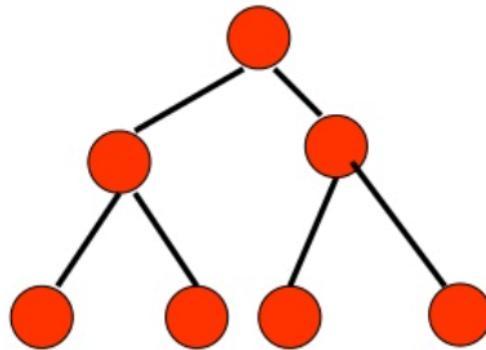
- Mean-shift for clustering



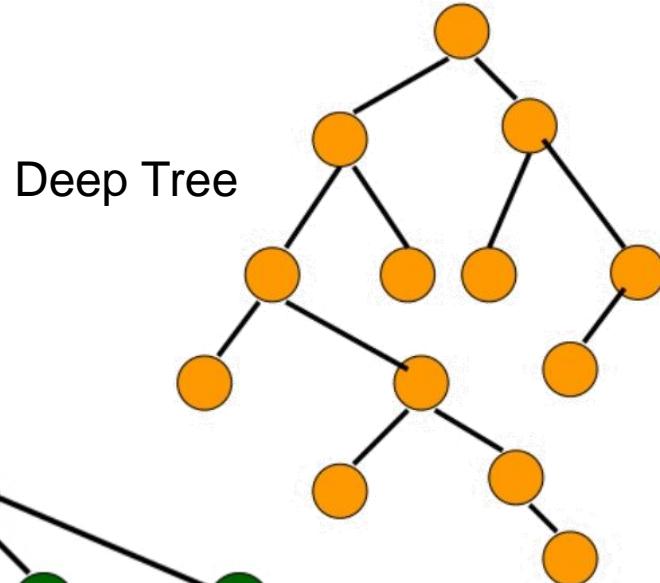
# Decision Tree

---

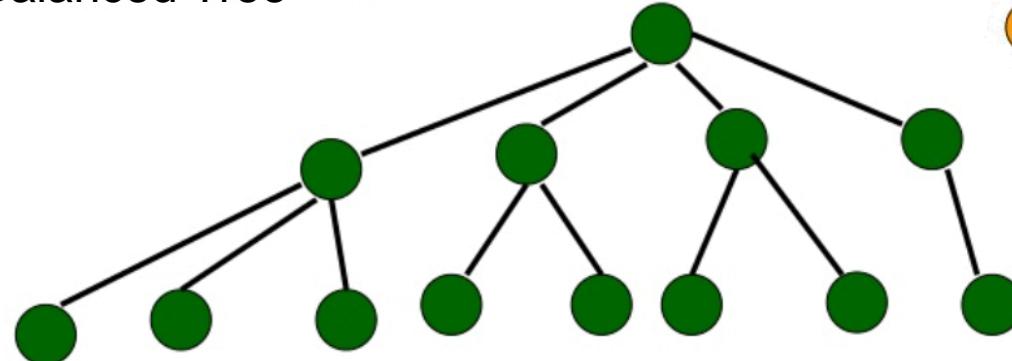
## ■ Types of Decision Tree:



Balanced Tree



Deep Tree

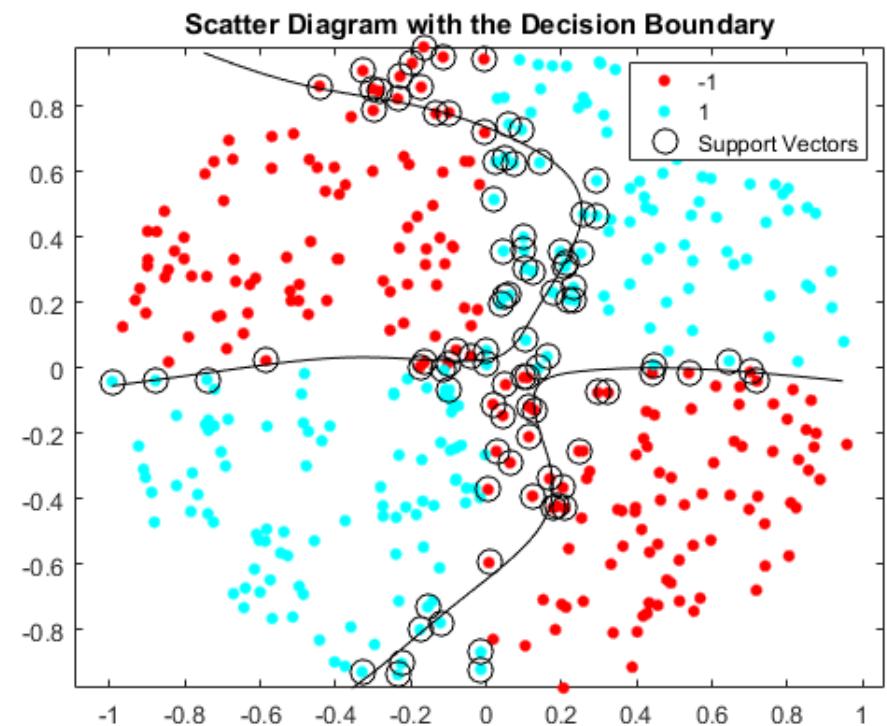
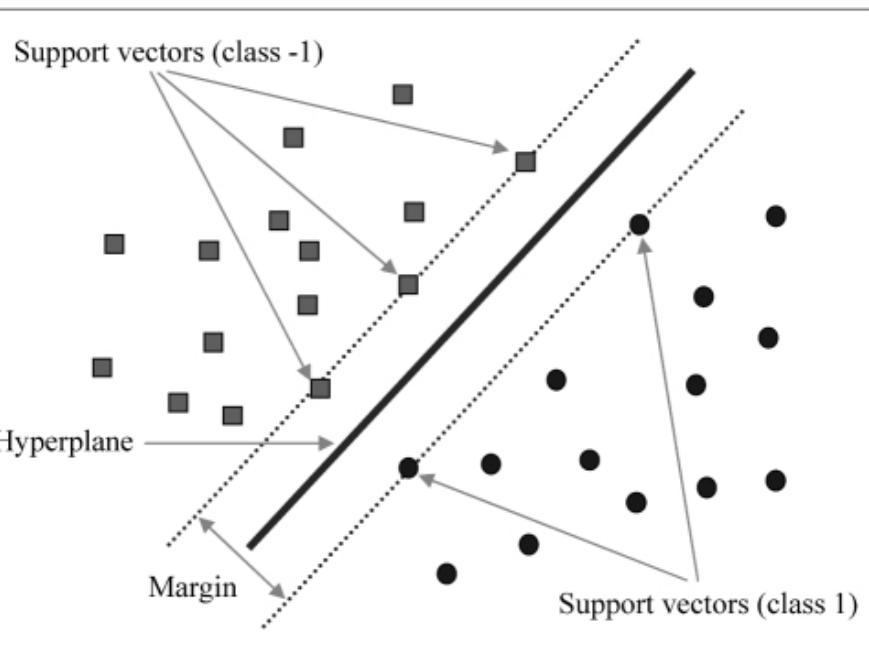


Bushy Tree

---

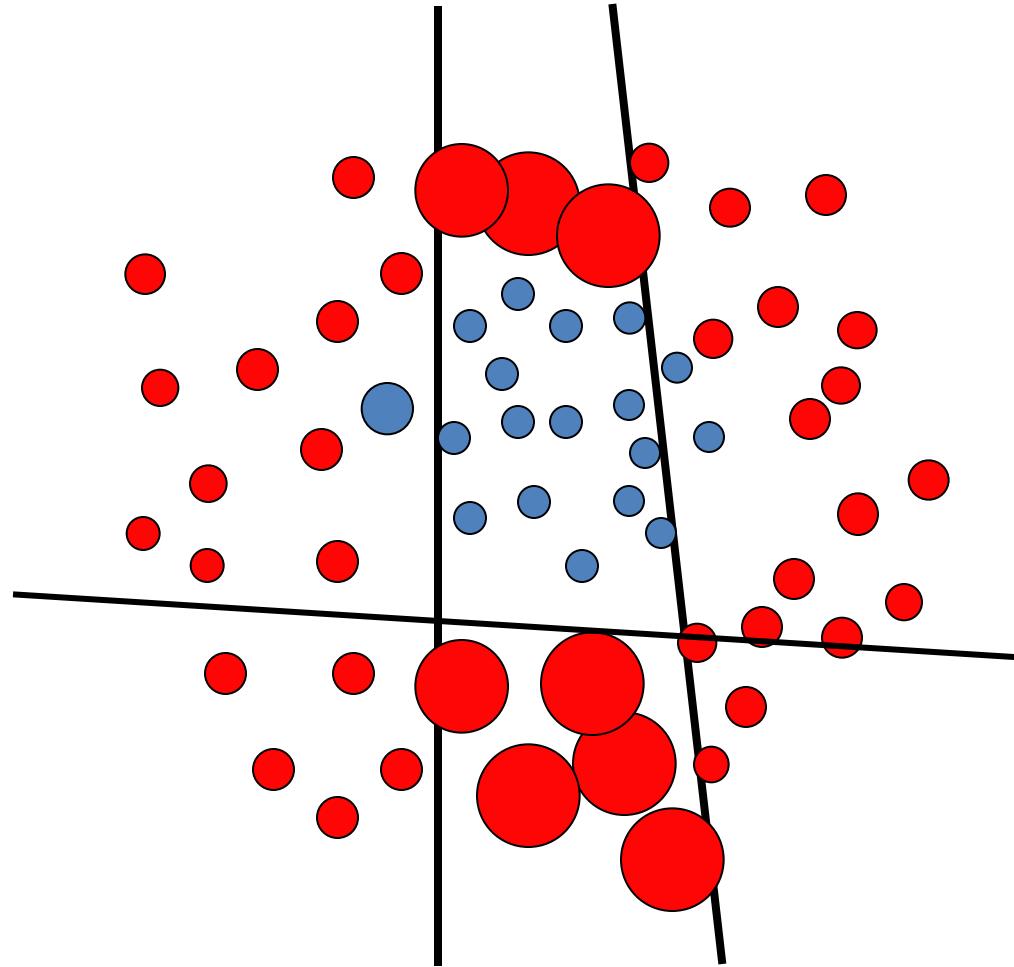
# SVM

- eg. Linear SVM:
$$\begin{aligned} & \arg \min_w \sum_{i=1}^n \|w\|^2 + C \sum_{i=1}^n \xi_i \\ & \text{s.t. } 1 - y_i x_i^T w \leq \xi_i \\ & \quad \xi_i \geq 0 \end{aligned}$$



# Boosting

---



Each data point has  
a class label:

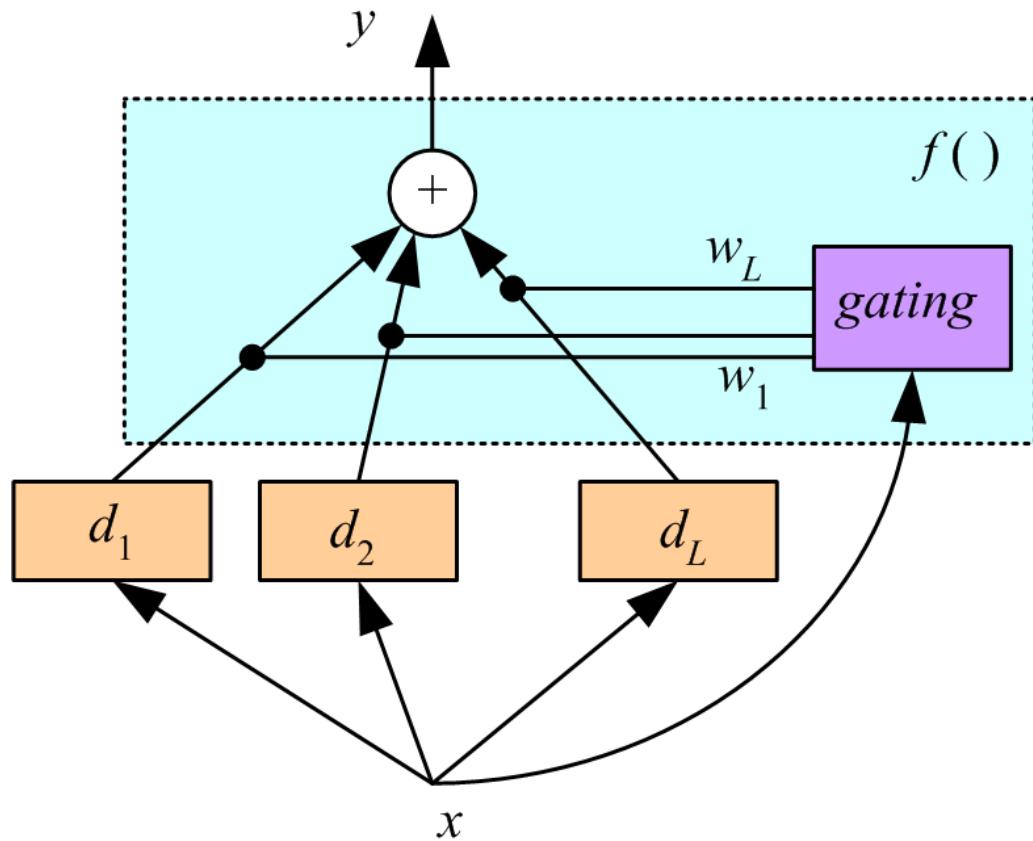
$$y_t = \begin{cases} +1 (\text{red circle}) \\ -1 (\text{blue circle}) \end{cases}$$

We update the weights:

$$w_t \leftarrow w_t \exp\{-y_t H_t\}$$

# Ensemble Learning

---



$$y = \sum_{j=1}^L w_j d_j$$

# Linear Statistical Learning

---

## ■ PCA

$$\begin{array}{c} \text{Data} \quad \text{Basis} \\ \searrow \quad \swarrow \\ Y = AX \quad \text{Coefficients} \\ A_i \perp A_j \end{array}$$

## ■ ICA

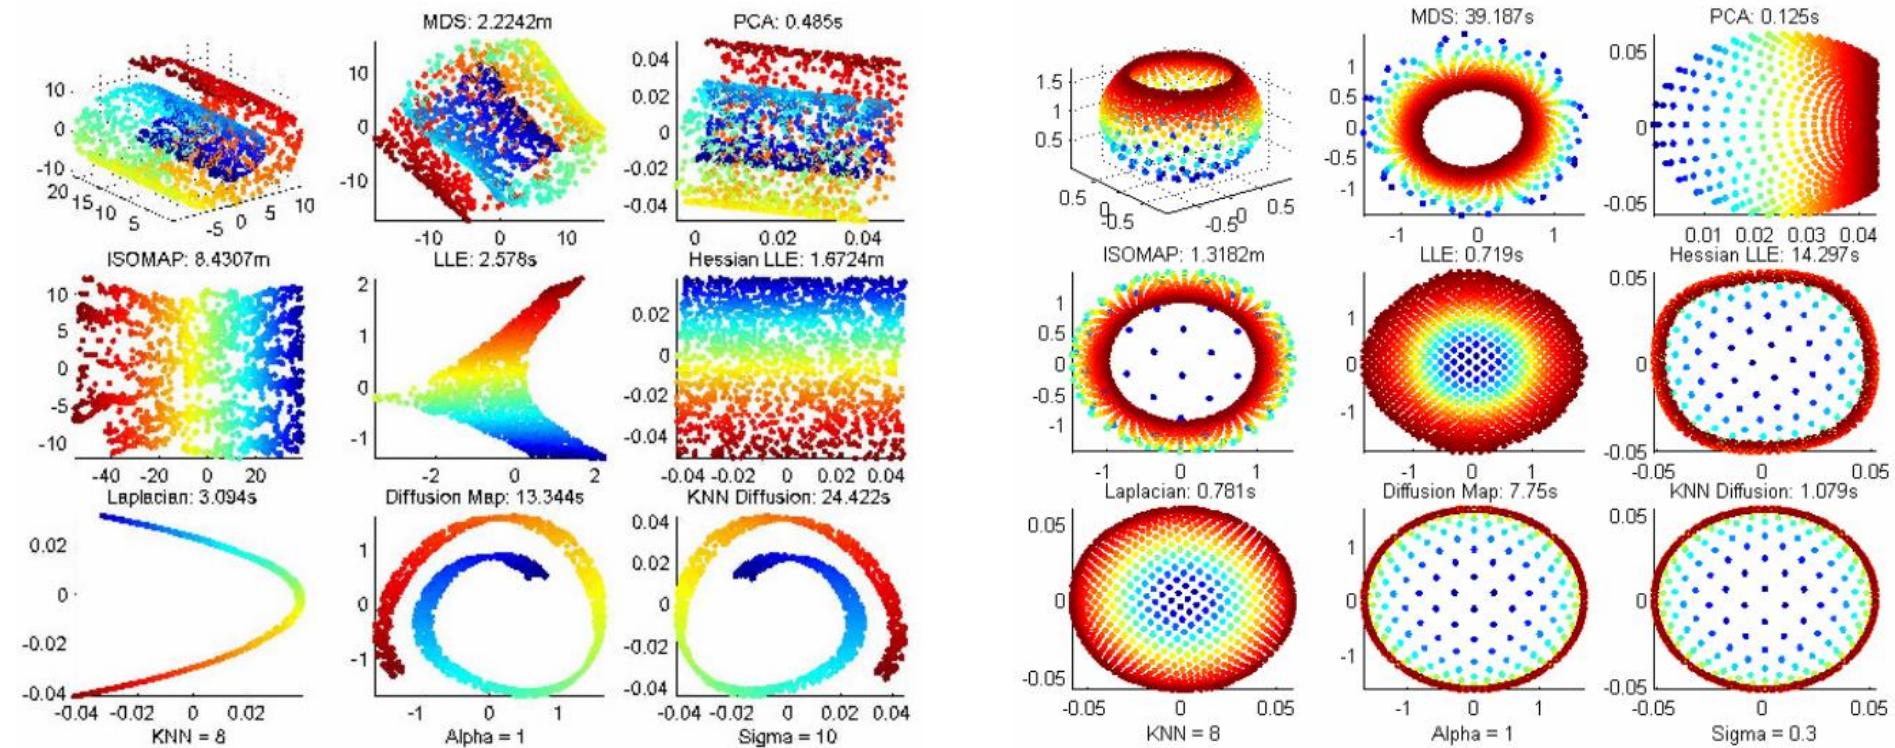
$$\begin{array}{c} \text{Data} \quad \text{Mixture Coefficients} \\ \searrow \quad \downarrow \\ Y = AX \quad \text{Components} \\ \max I(X) \end{array}$$

## ■ NMF

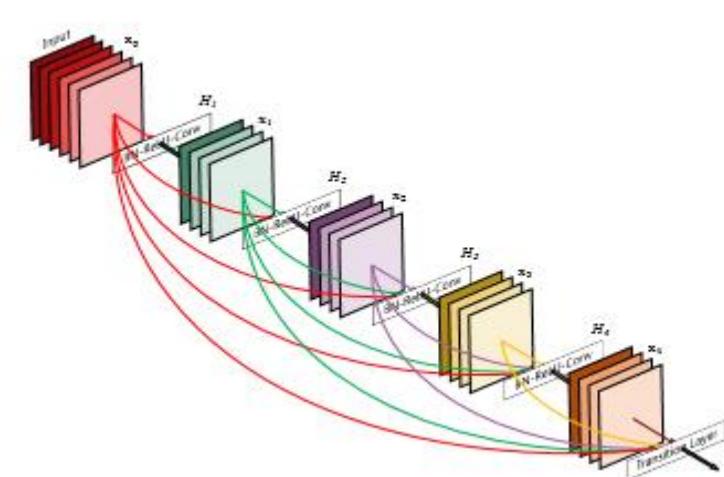
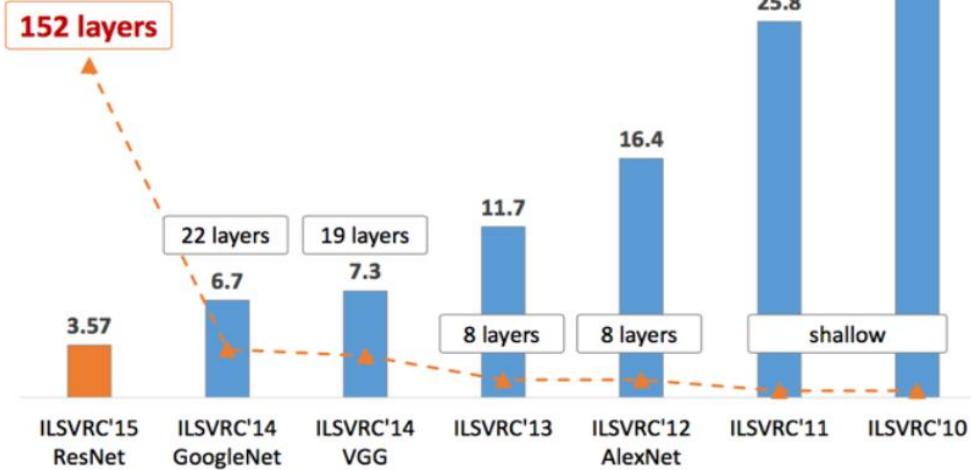
$$\begin{array}{c} \text{Data} \quad \text{Basis} \\ \searrow \quad \swarrow \\ Y = AX \quad \text{Coefficients} \\ A, X > 0 \end{array}$$

# Nonlinear Statistical Learning

## ■ Manifold learning



# Deep Neural Networks

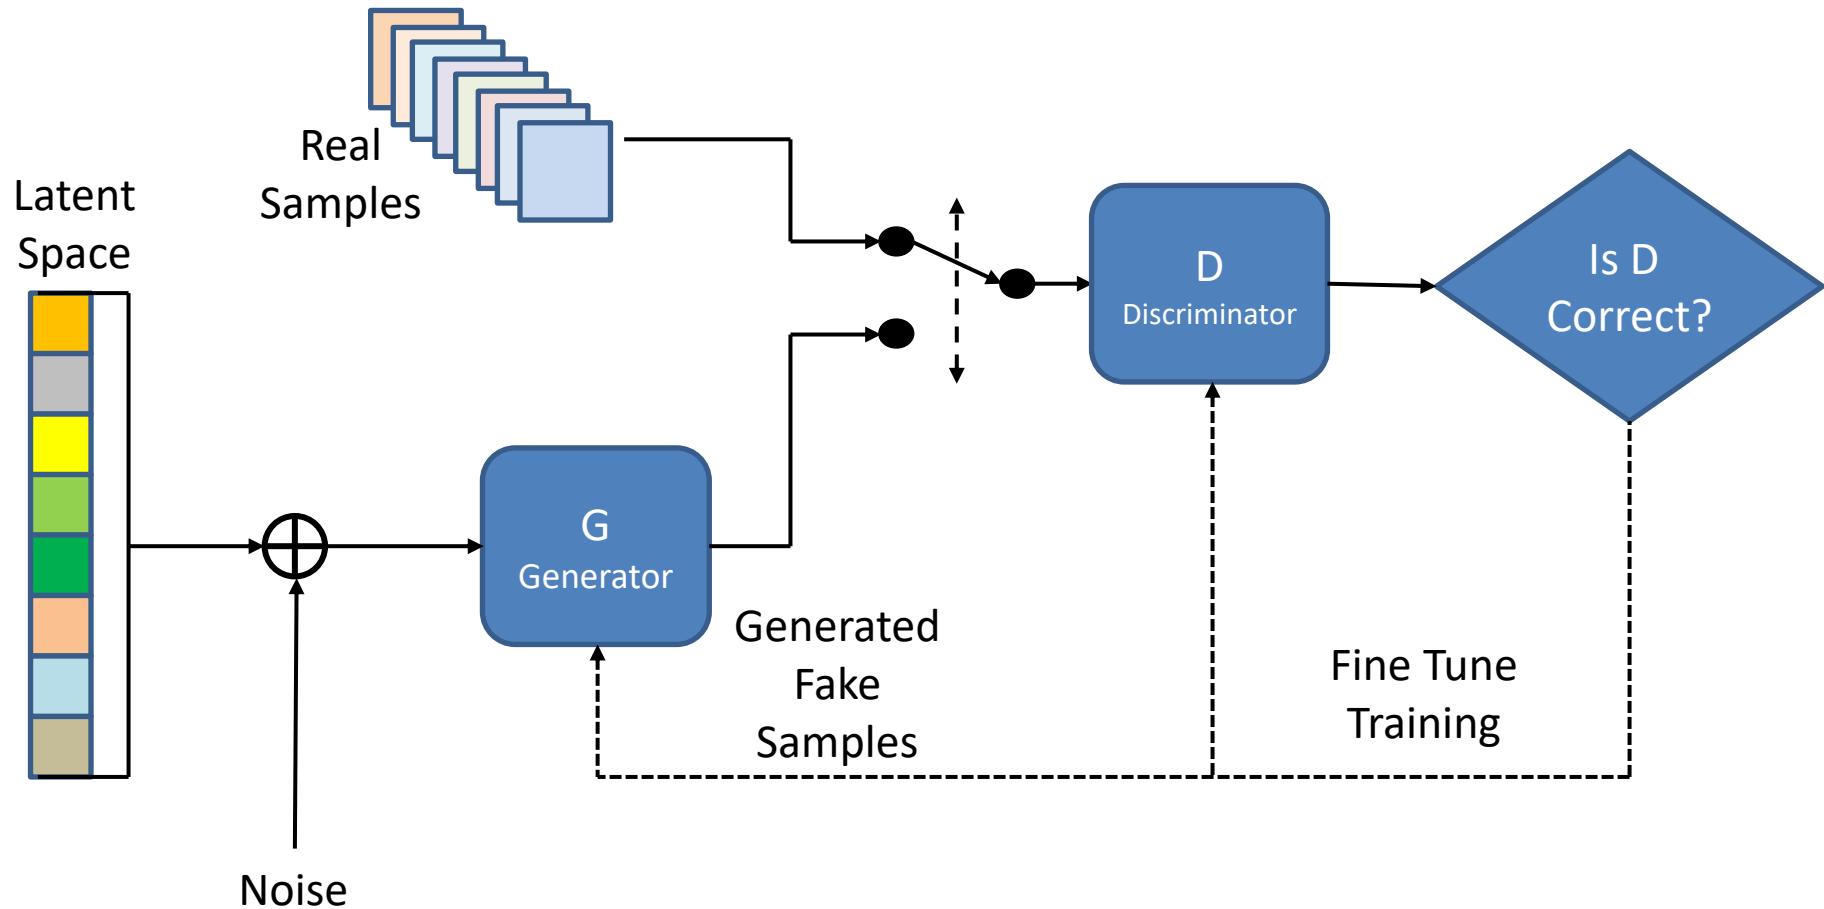


- Task: recognition
- Dataset: ILSVRC

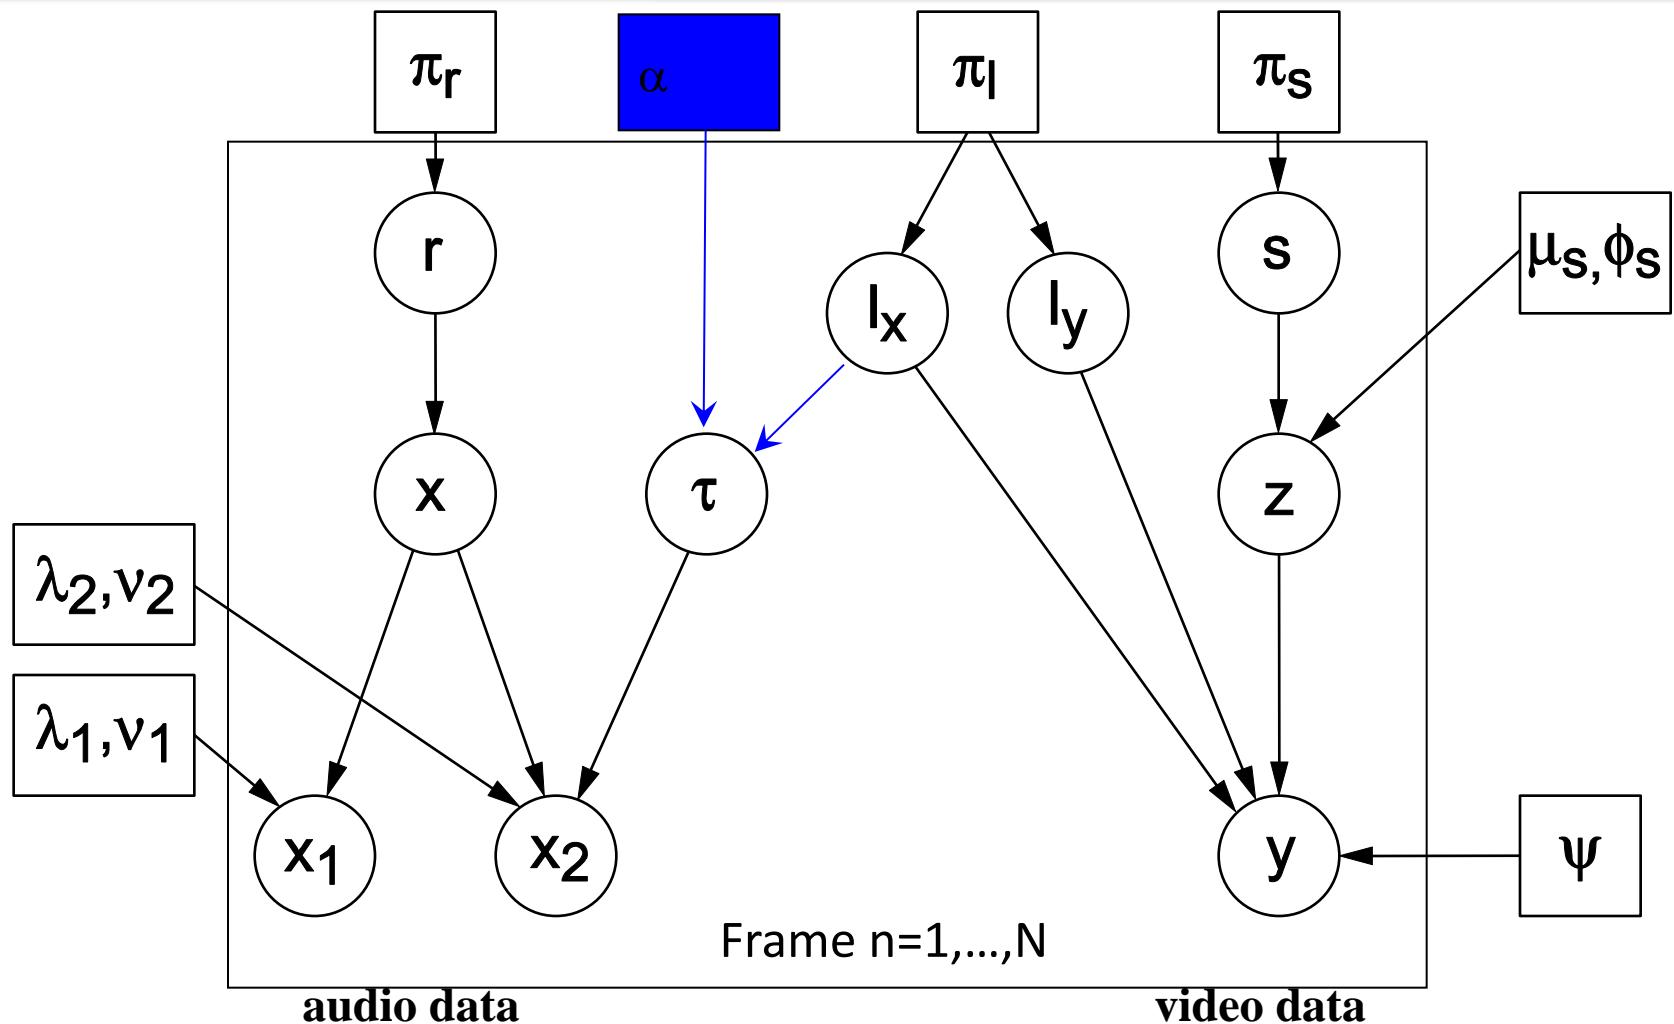
■ Huang G, Liu Z, Weinberger K Q, et al.  
Densely connected convolutional networks[J]. arXiv preprint  
arXiv:1608.06993, 2016.

# Generative Adversarial Networks

---



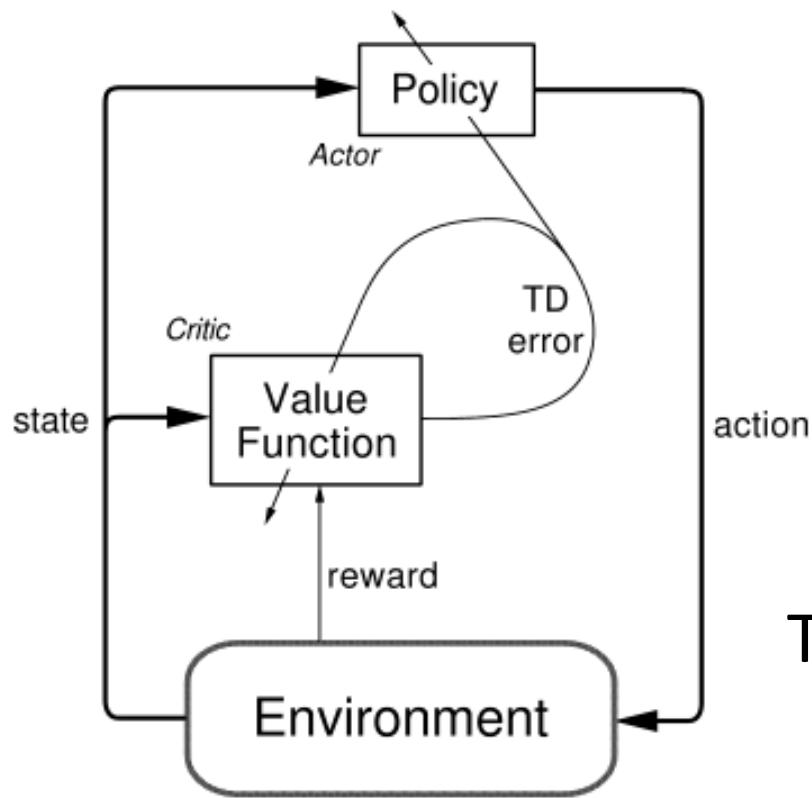
# Bayesian Networks



# Reinforcement Learning

---

- State, action, and Reward



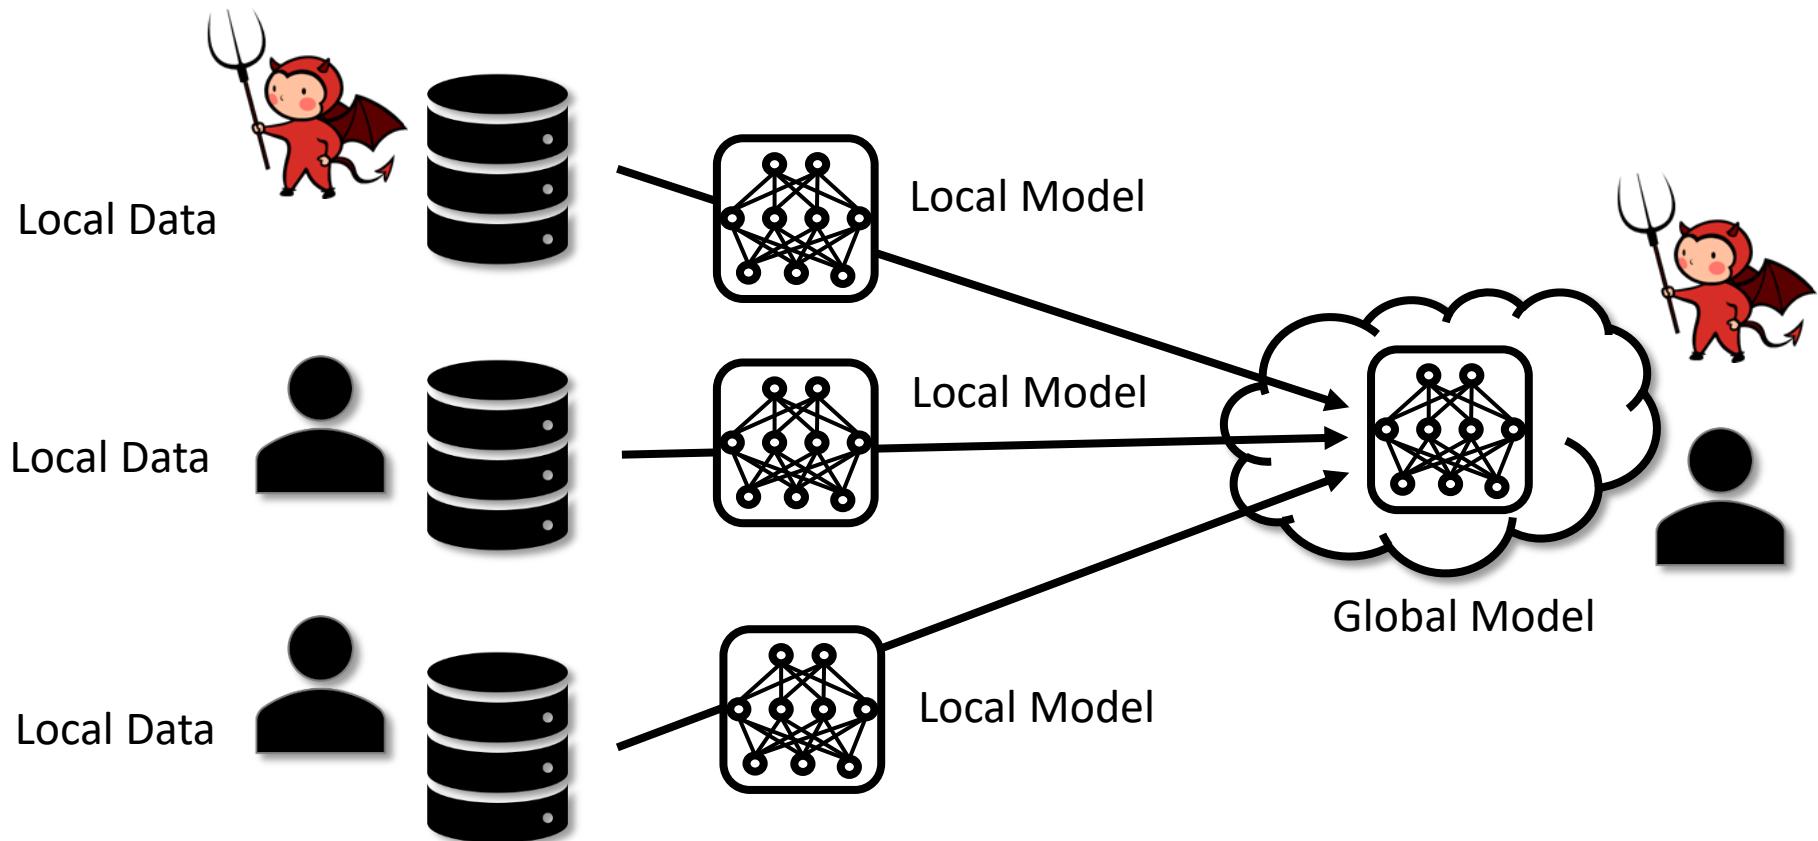
Update: Policy Function  
Value Function

TD Error: Temporal Difference  
between Real Reward  
and Estimated Reward

# Federated Learning

---

- Collaborative Learning



# Outlines

---

- Framework
  - Problem Statement
  - Related Areas
  - History
  - Datasets and Learning Models
  - Optimization Methods
  - Algorithms
  - Examples
-

# AlphaGo

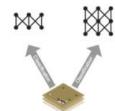
AlphaGo vs. 李世石



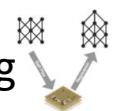
AlphaGo vs. 柯洁



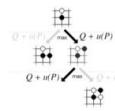
Supervised learning



Reinforcement learning



Priori knowledge



Better network **structure**

Enhance the role of **Reinforcement learning**

# Autonomous Driving

2014 Google



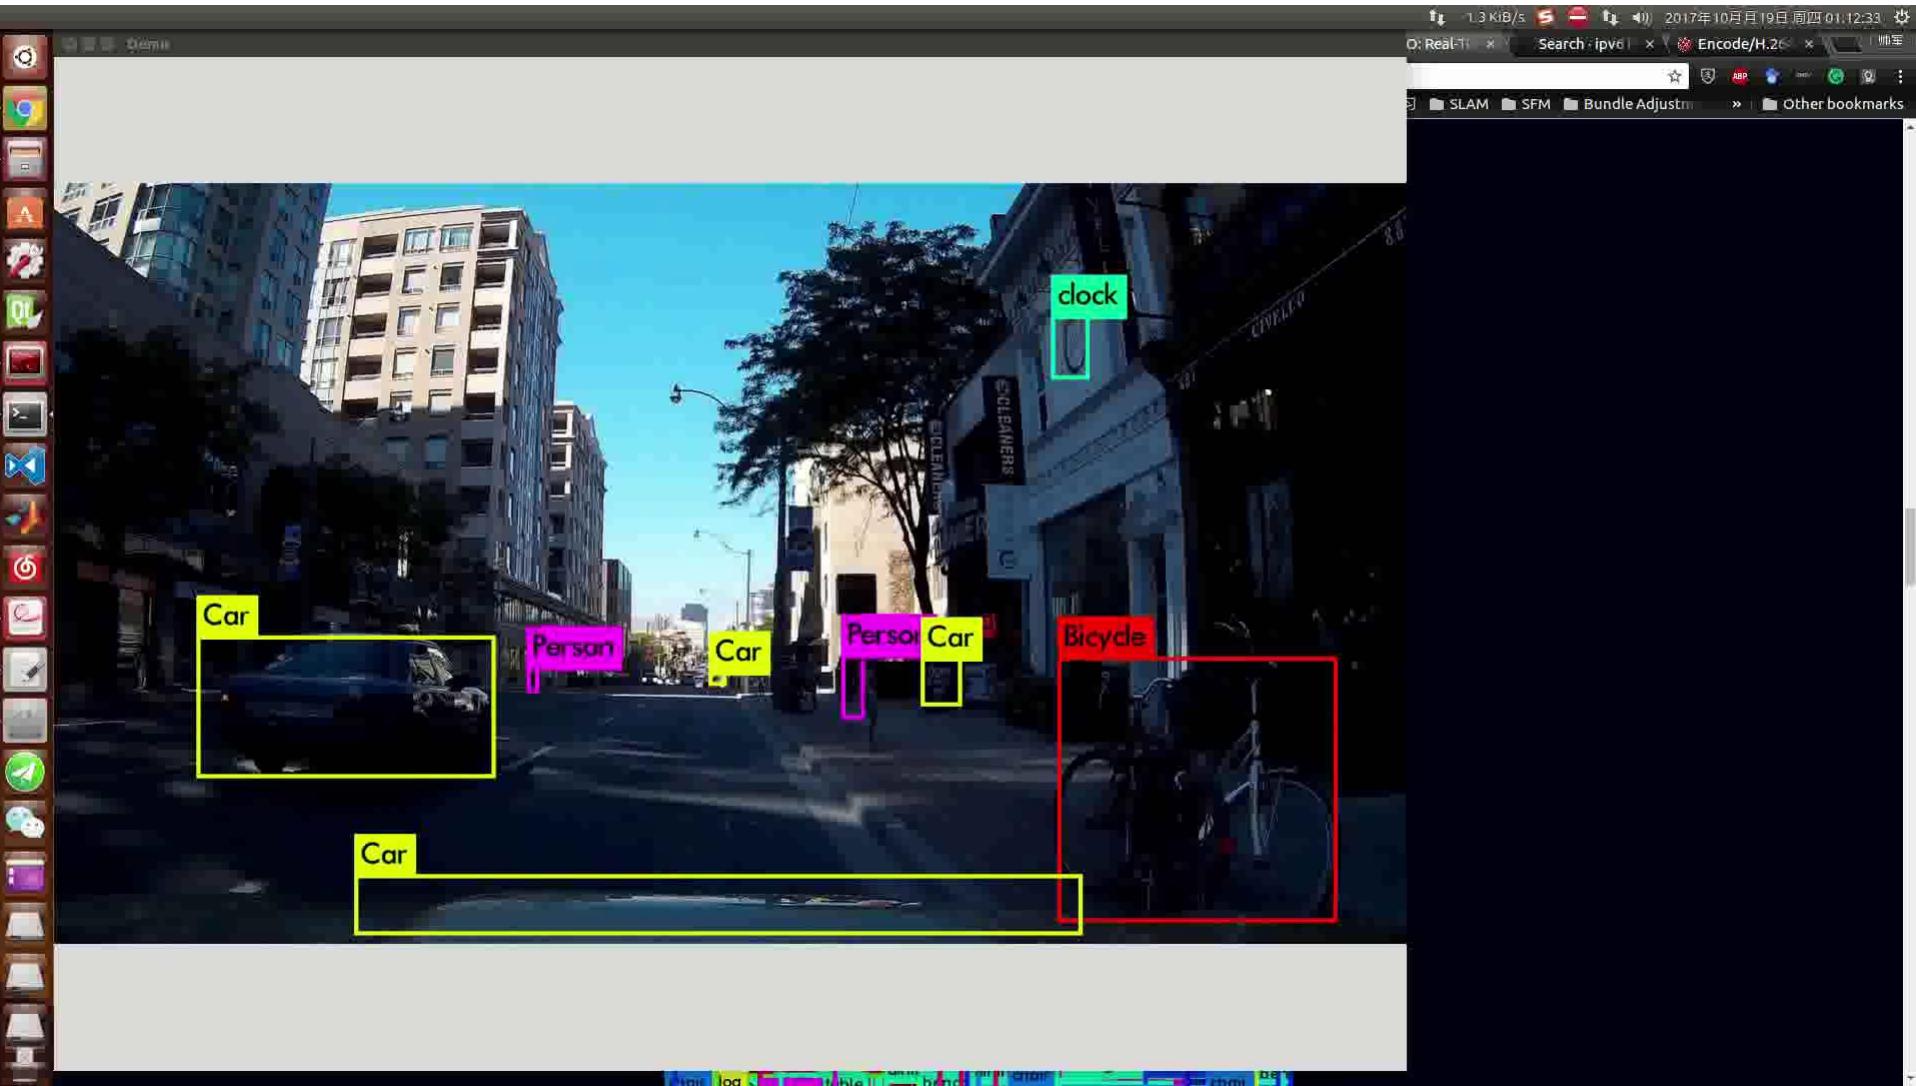
2016 Tesla



2017 apollo



# Object Detection-YOLOv3



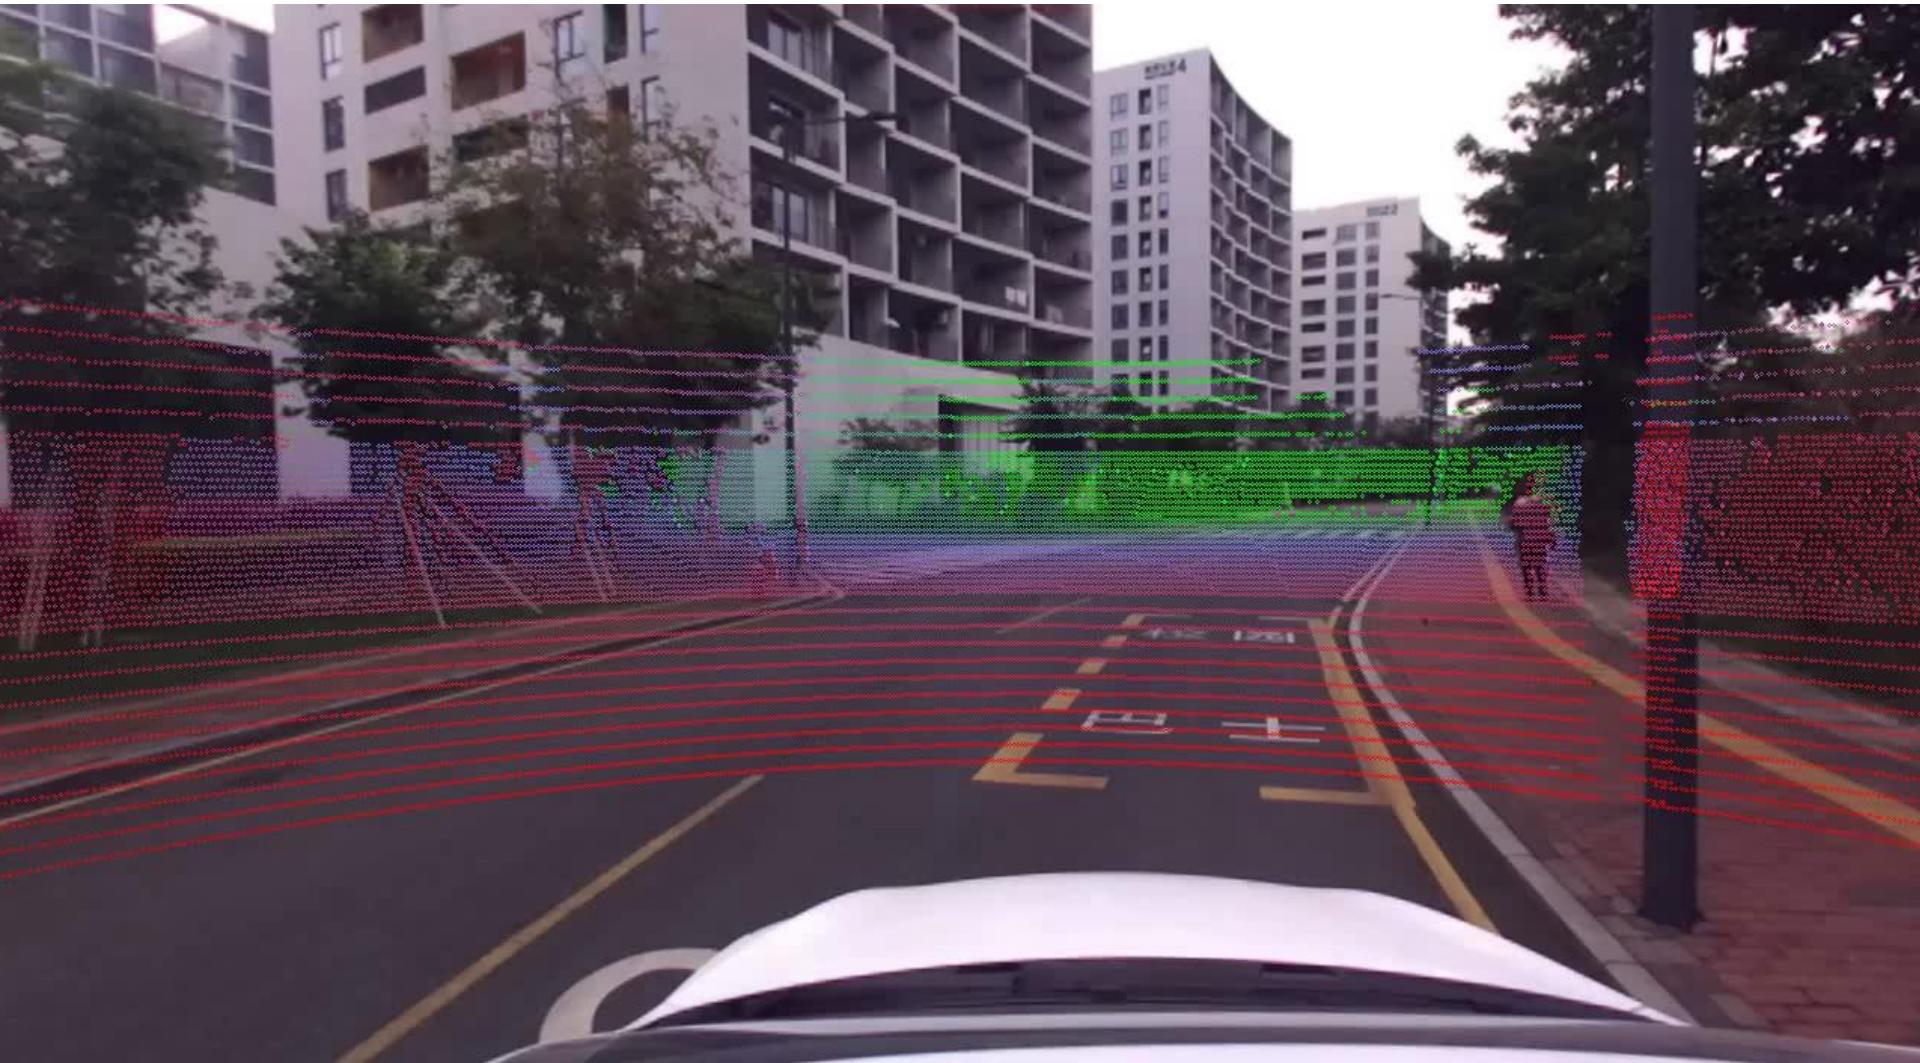
# Instance Segmentation



# Lane & Sign Detection

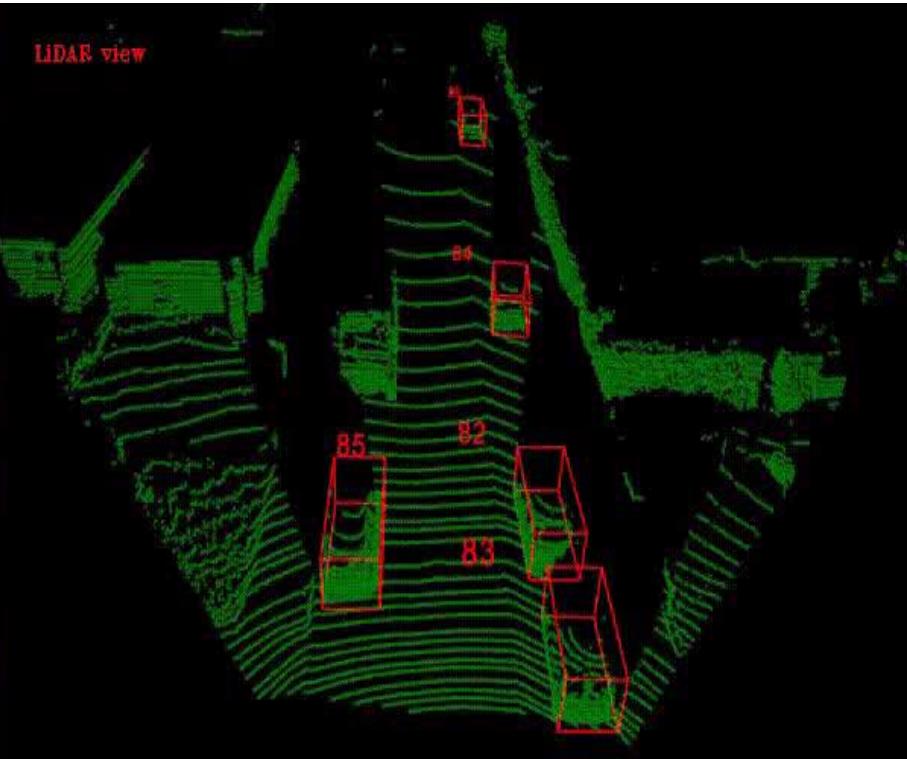
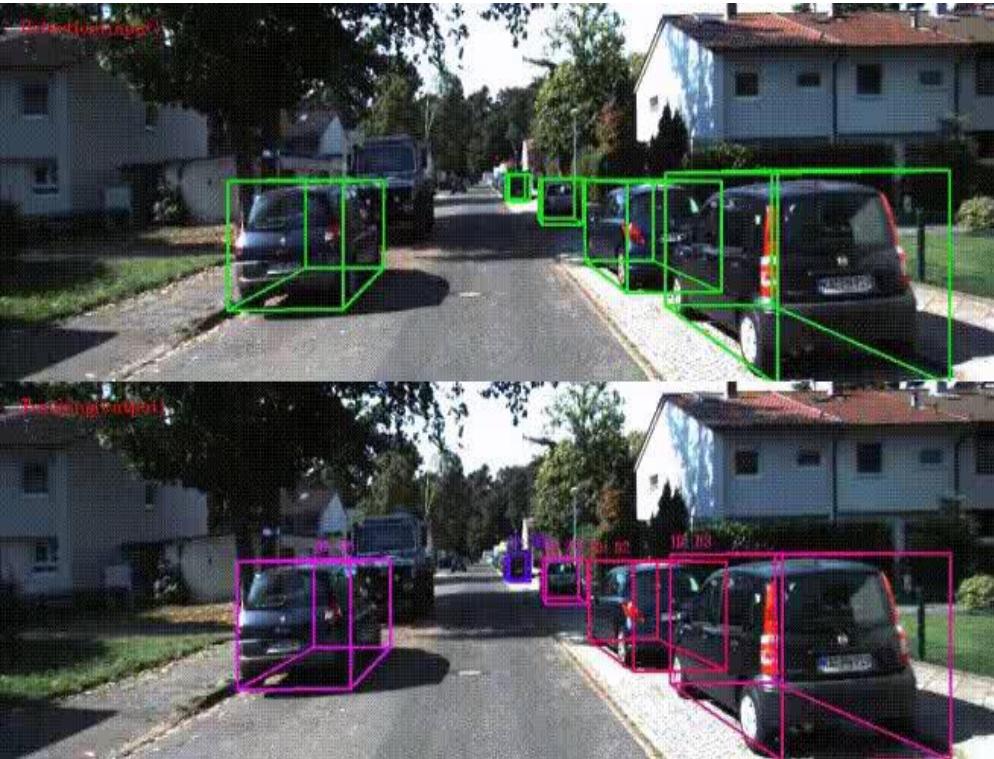


# 2D-3D Fusion



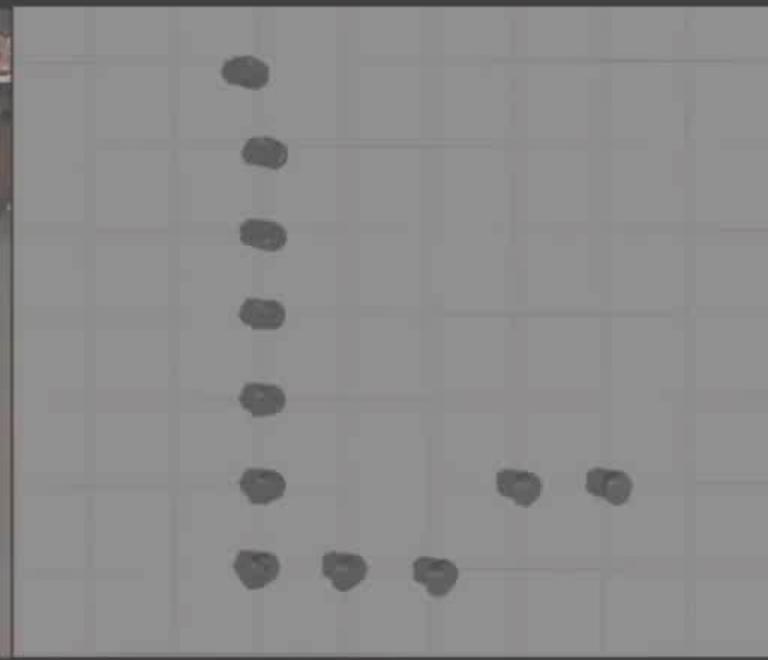
# 2D-3D Fusion for Tracking

---



# Reinforcement Learning for Navigation

2.5X



S U S T E C H

# Federated Learning for Map Fusion

Scenario Generation In Carla

# Datasets and Benchmark Metrics

---

## ■ Datasets

- KITTI
- Udacity
- Waymo
- NuScences
- CityScapes
- ApolloScape
- SUSTech Scape



## ■ Metrics

- Correctness
  - Robustness
-

# Our Dataset

**SUSTechscape**

*An open datasets for autonomous driving*

[Home](#)

[Datasets](#)

[Download](#)

[Benchmarks](#)

[Simulator](#)

[Submit results](#)

[Collaborators](#)

[Contact](#)

# SUSTech Scape

Open Datasets for Autonomous Public Transportation with Smart Samples and Cyber-Physical Benchmarks

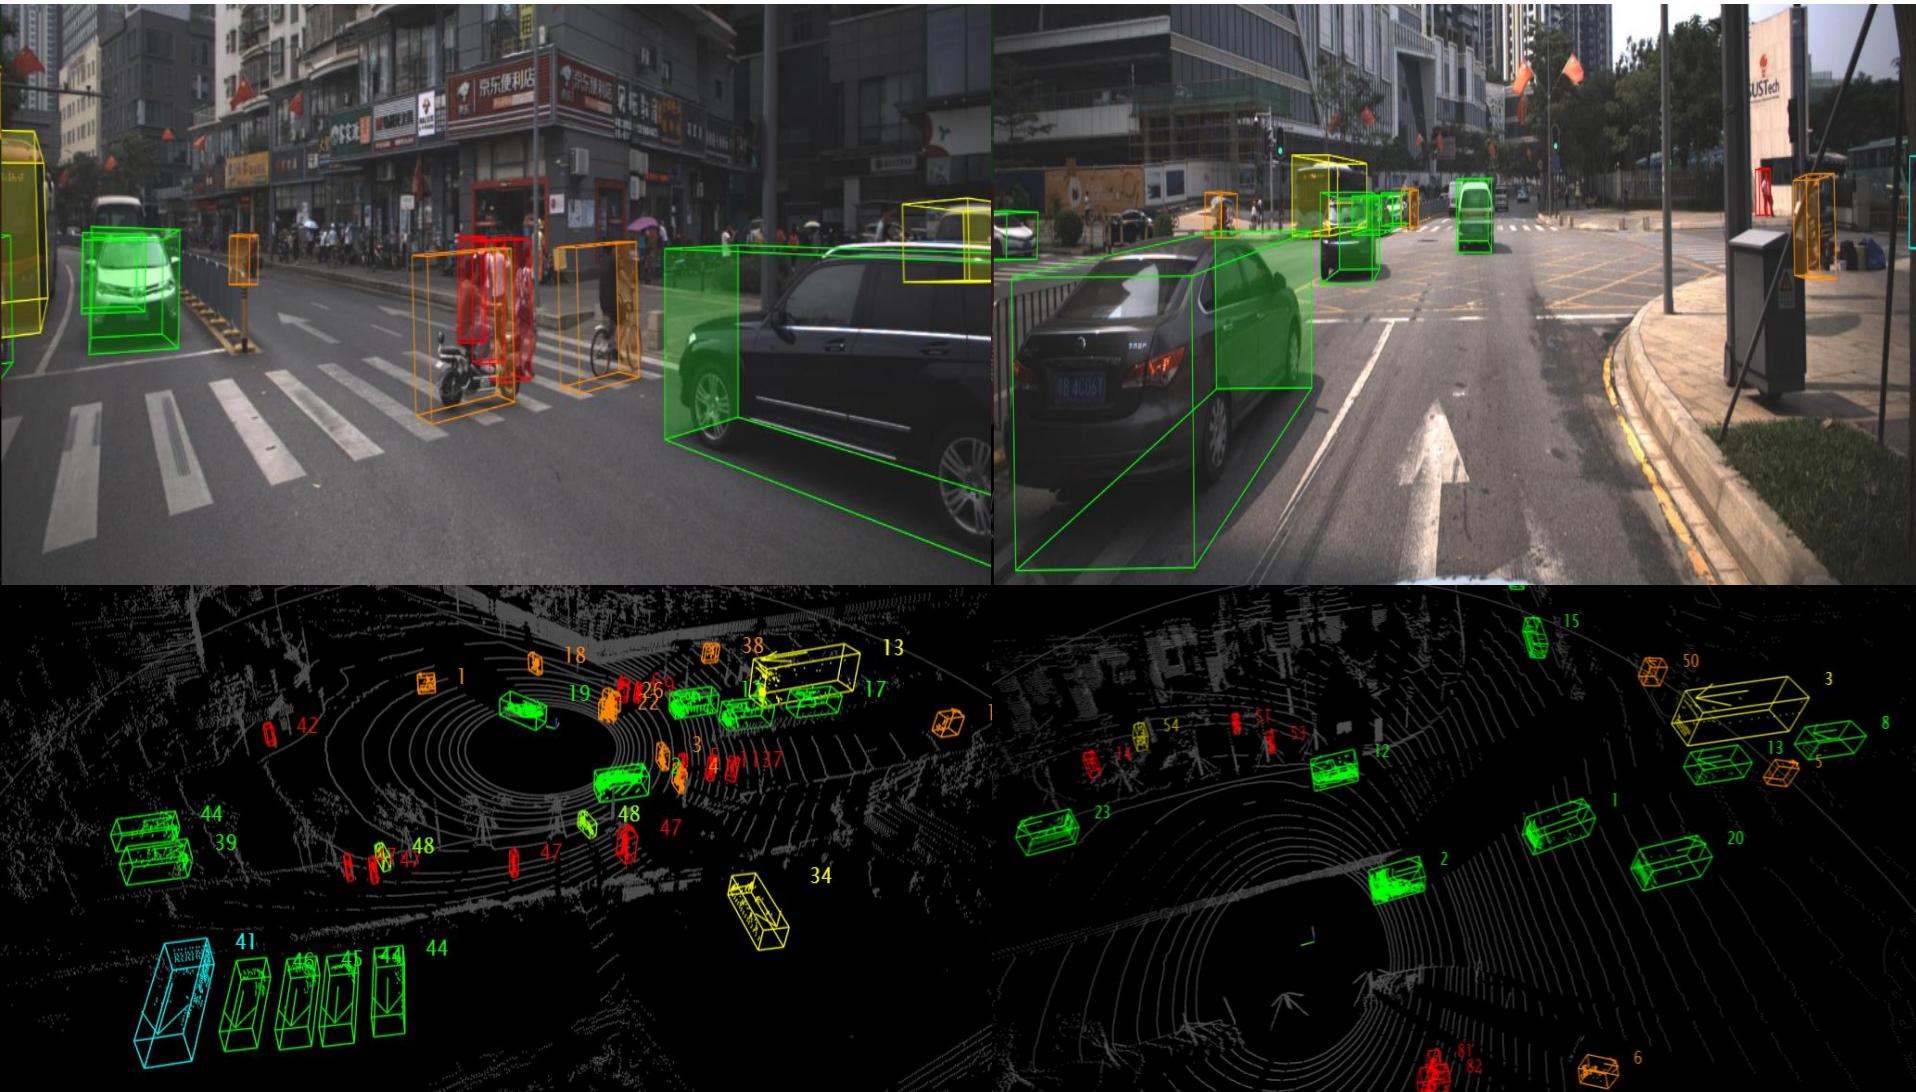


# Datasets Comparison

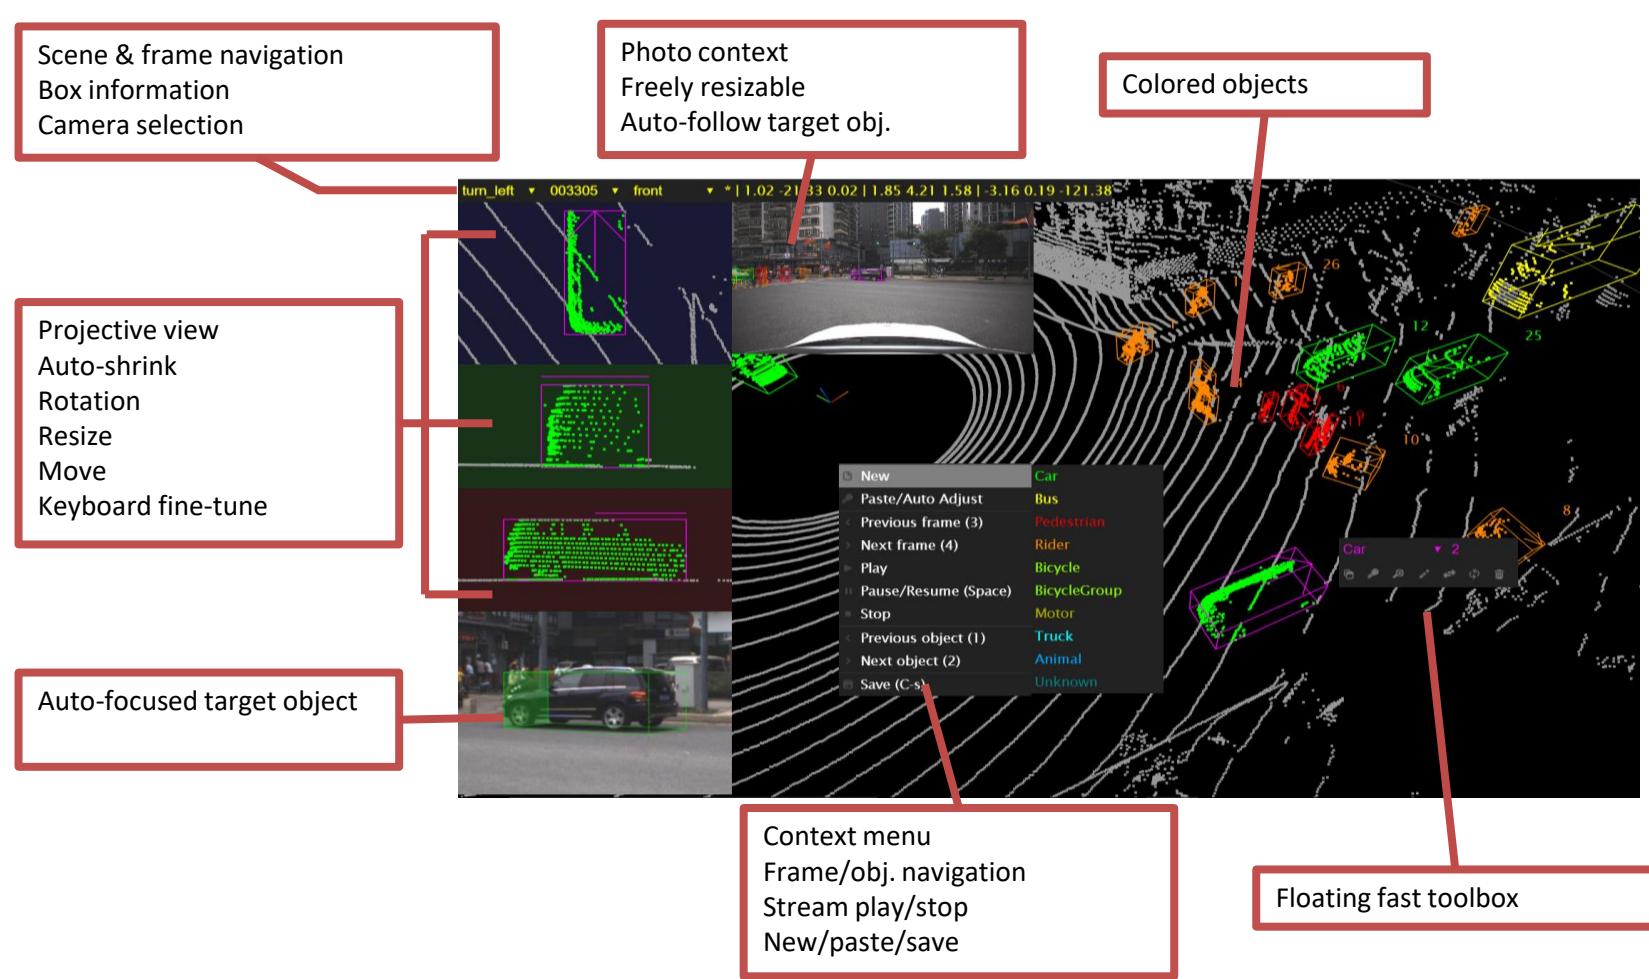
Dataset	Sensors	Location Accuracy	Scene Diversity	Annotation				Driver	
				2D	3D	Video	Lane	Behav	Physi
KITTI	Cameras (6), Lidar (1)	cm	Regions Day time	box-15k, pixel-400	Box	no	no	no	no
Oxford RobotCar	Camera, Lidar	unknow	1 City, weather	no	no	no	no	no	no
TSD-max	Camera	m	Regions	box-10k	no	no	2D	no	no
CityScape	Camera	unknow	50 Cities, Season, weather, D&N	pixel-25k	no	no	no	no	no
Udacity	Camera, Lidar	unknow	1 City, sunny and overcast	box-25k	no	no	2D	no	no
Mapillary	Camera	m	Cities, season, weather, D&N	pixel-25k	no	no	2D, 2 classes	no	no
TorontoCity <sup>1</sup>	Camera, Lidar	cm	1 City	pixel	point	no	no	no	no
BDD100K	Camera, Lidar	m	4 Regions, weather, D&N	box-100k, pixel-10k	no	no	2D, 8 classes	no	no
ApolloScape	Cameras (6) Lidars (2) Cars (4)	cm	4 Regions, weather, D&N	pixel-140k resolution 2K	point	yes	2D, 3D, 28 classes	no	no
SUSTech Scape	Camera Lidars (3) M-spectral Cameras (2) Buses (4) Taxis (4)	(6)	cm	1~5 Cities, weather, day and night	pixel-200k resolution <4K	point	yes	2D, 3D >28 classes	yes

Benchmark Metrics					
Evaluation Metrics			Methods	Measures	Subjects
Scene Processing	Scene flow	stereo disparity outliers [%] optical flow outliers [%] scene flow outlier [%]	Image, 3D Points Cloud Based Experiments	Optical and scene flow estimation errors, 3D reconstruction errors compared with ground truth	Optical flow estimation, scene flow estimation, 2D-3D reconstruction algorithms
	Depth	Scale invariant error [%] Relative squared error [%] Absolute squared error [%] Root mean squared error [%]			
Mapping & Localization Accuracy	Mapping	iRMSE [1/km] iMAE [1/km] RMSE [mm] MAE [mm]	Odometry (IMU, Camera) Readings and 3D Points Cloud Based Experiments	Environment mapping and vehicle localization under various scenarios and conditions	SLAM, key frame selection, sensor fusion, optimization algorithms
	Odometry	Translation errors [%] Rotation errors [deg/m]			
Energy Efficiency	Urban Area	Battery capacity [kW/h] Range [kM] Complete weight [t]	Testing Experiments	Current and power consumption of devices	Power modes, battery management and Scheduling algorithms
	Suburb Area				
	Hybrid Area				
System Scalability	Task completion time [h] Task completion distance [km] MTCA [km] MTDA [h]	Simulation Experiments in Different Scales	Predicted system performances in different scales of the vehicle fleet	Communication protocols, group formation control, collaborative planning Algorithms	
System Reliability & Robustness	Scenarios Overlap Rate[%] One Pass Rate[%] Defect resolution rate[%] Pressure running time[%]				
		Simulation Experiments under Complex Scenarios	Sharp turns, degraded marks, faulty signals, reckless driving, sensor failures	Sensor perception, motion planning, decision & control algorithms	

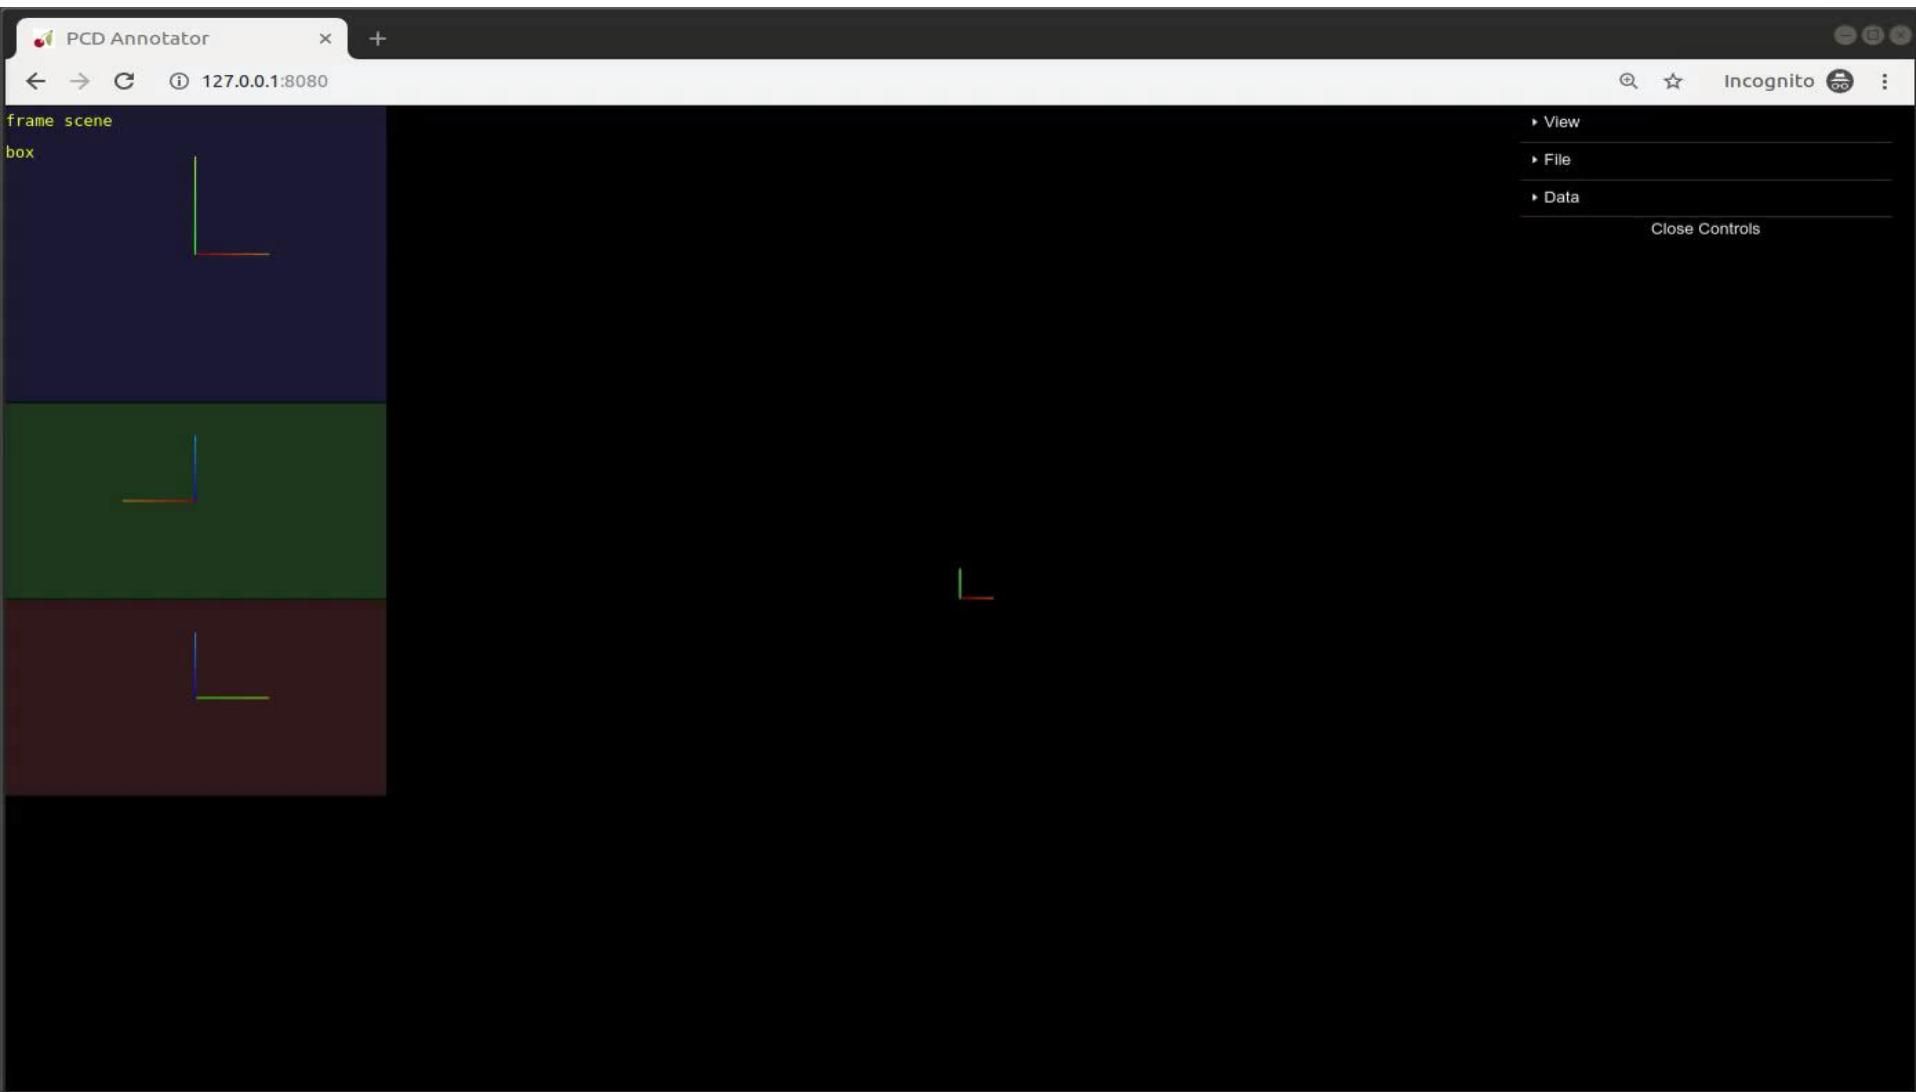
# Dataset Annotation



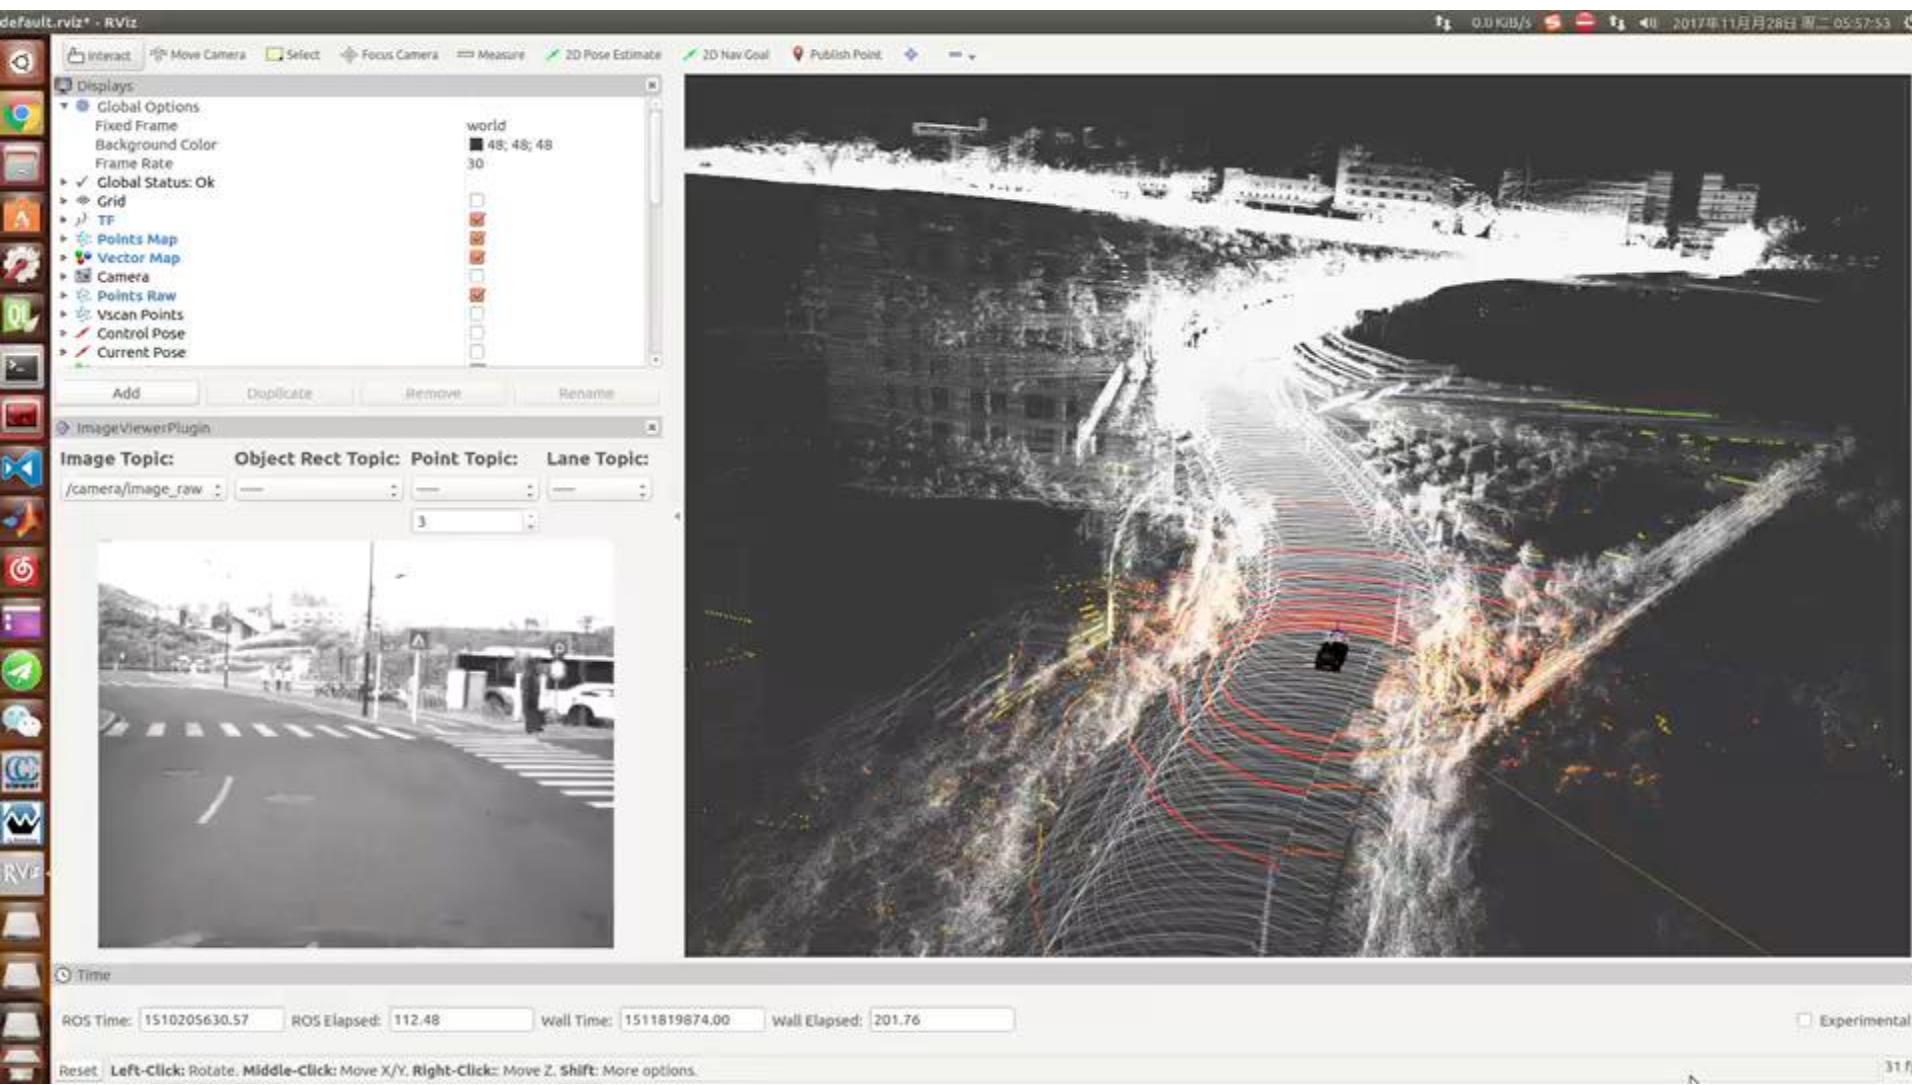
# Dataset Annotation Tools



# SUSTech POINTS

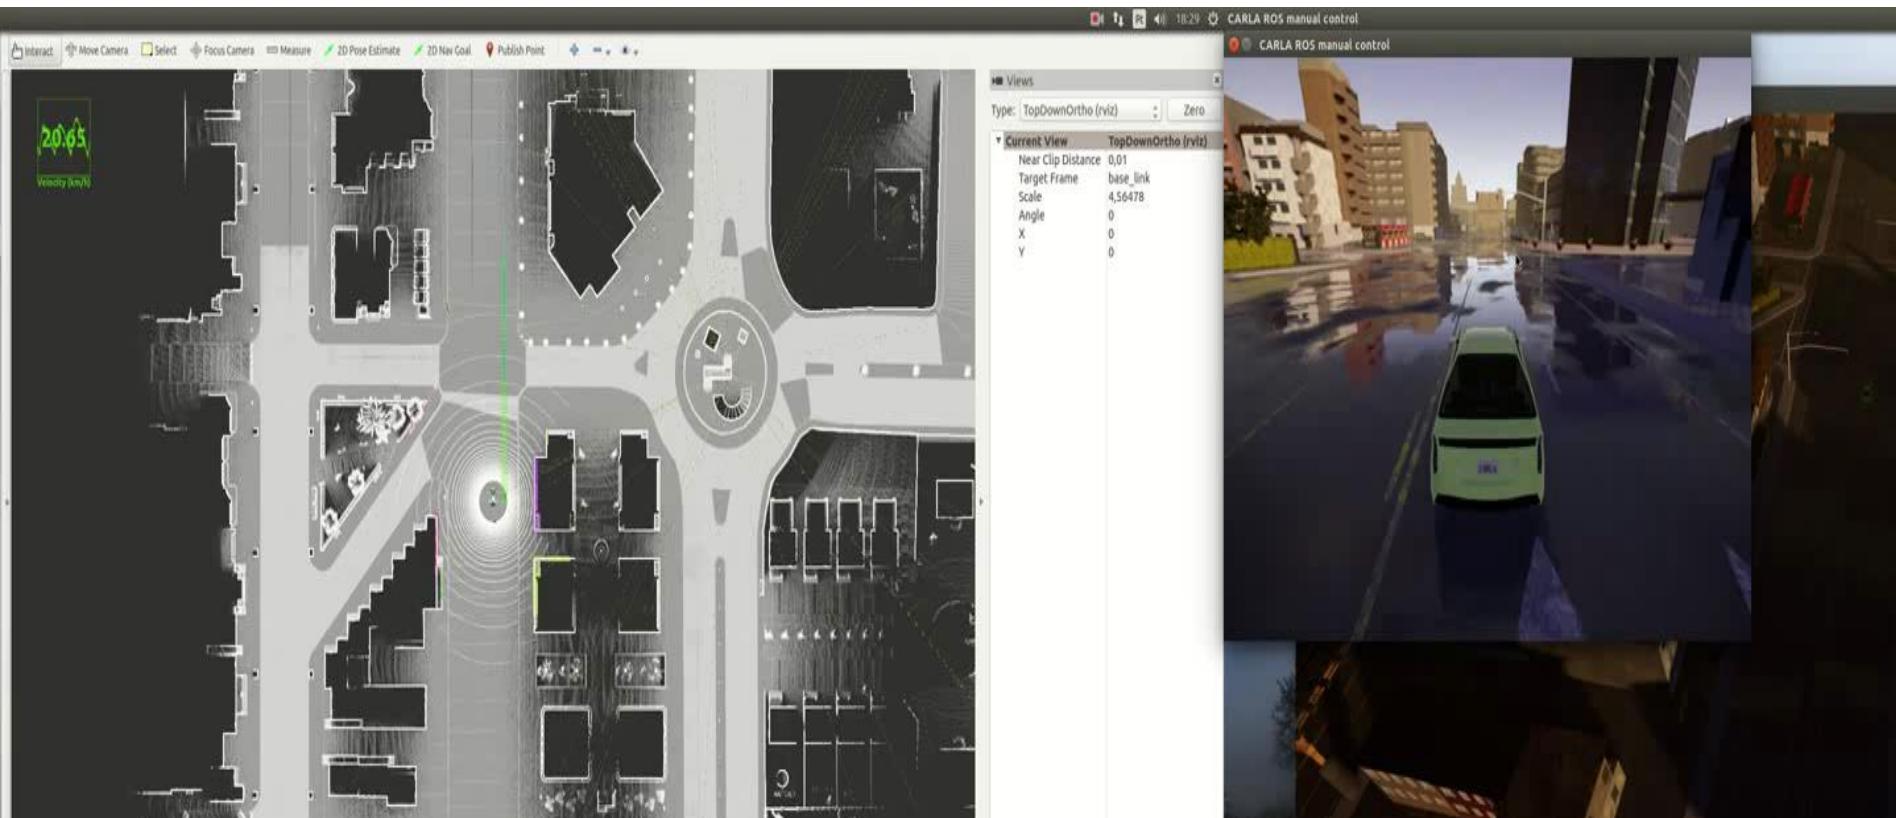


# Play-Back: Autoware



# Simulators: CARLA+Autoware

---



# More Reading and Multimedia Materials

---

**Books:** 《人类简史》 《奇点将近》 《终极算法》  
《人工智能时代》 《2050》 《情感机器》  
《数学之美》

**Movies:** “Blade Runner” “AI” “Prometheus”  
“Covenant” “Ex Machina” “She”  
“2001: Space Odyssey” “The Matrix”  
“I, Robot” “Bicentennial Man”  
“Terminator”

**TV Series:** “West World” “Humans” “Black Mirrors”

---

# More Course Links

---

**Stanford Machine Learning:**

<https://see.stanford.edu/Course/CS229/47>

**MIT Machine Learning:** <https://ocw.mit.edu/courses/electrical-engineering-and-computer-science/6-867-machine-learning-fall-2006/index.htm>

**Stanford CNN for Vision:** <http://cs231n.stanford.edu>

**Stanford Deep Learning:** <http://cs230.stanford.edu/syllabus.html>

**MIT Deep Learning:** <http://introtodeeplearning.com/>

---